

TET15-P1

A Three-Field Hu–Washizu Tetrahedron

Research note

Alejandro Mota

October 4, 2026

Abstract

TET15-P1 is a tetrahedral finite element for the large-deformation plasticity of nearly incompressible materials: materials whose bulk modulus κ , the resistance to a change of volume, is much larger than their shear modulus μ , the resistance to a change of shape, as in rubber and in metals during plastic flow. Three fields — the motion, a scalar volumetric strain and a pressure — are coupled through a Hu–Washizu functional, a functional in which the two auxiliary fields are independent of the motion and tied to it by a constraint term. With both auxiliary fields in one discrete space that is discontinuous between elements, the functional reduces exactly to a functional of the motion alone in which the volumetric strain $\theta(J)$, a function of the Jacobian J of the motion, is replaced element by element by its L^2 projection onto that space (mean dilatation), and the auxiliary fields are eliminated element by element. The deviatoric (shape-changing) response is evaluated at the quadrature points, so the constitutive law and its return map, the algorithm that projects a trial stress onto the yield surface, run unchanged, and no quantity other than θ is projected. In its general form the element calls the material at a modified deformation gradient whose volume is the projected one, and needs nothing of the material beyond its stress and tangent; when the stored energy splits exactly into a volumetric part quadratic in θ and an isochoric (volume-preserving) part, the pressure equals κ times the projected strain at every point and the element needs one projection and one material call per quadrature point. Both forms are implemented; the general one takes one pass more over the quadrature points and, with the J_2 model, 1.4 to 2.1 times the split form’s time per explicit step on six devices. The choice of $\theta(J)$ is a parameter of the formulation, and two instances of the split form are developed, $\theta = \log J$ with the Hencky material and $\theta = J - 1$ with the J_2 plasticity model of Simo and Hughes, the instance implemented. The discrete spaces are the three-dimensional Crouzeix–Raviart pair: the quadratic tetrahedron enriched with four cubic face bubbles and one quartic interior bubble, polynomials that vanish on a prescribed part of the element boundary, carried by a nodal basis on fifteen nodes (the TET15 displacement element), over a pressure that is linear in the reference coordinates of each element and discontinuous between elements; the name of the element states the two spaces, fifteen nodes for the displacement and a linear (P_1) pressure. Volumetric locking, the excess stiffness of a discretization as κ/μ grows, and spurious soft modes, discrete displacements of vanishing energy, are failures of the two hypotheses of Brezzi’s theorem for the pair of discrete spaces: the inf-sup condition, that every discrete pressure is resisted by some discrete displacement with a constant independent of the mesh, and coercivity of the deviatoric form on the discretely isochoric subspace. The inf-sup constant of the pair is measured over a mesh sequence at 0.2968, against 0.2216 for the Taylor–Hood pair (continuous quadratic displacement over continuous linear pressure), whose stability is

established; of seven element-local pairs two are stable and five decay. Under a plastic tangent that removes the shear stiffness along the flow direction, the coercivity constant cannot separate pairs, and what separates them is the dilatation carried by their softest modes: 0.20 to 0.27 without decay under refinement for elements with a constant pressure, the composite tetrahedron of Albany-LCM with its volume-averaged J among them, against a decay as the square of the mesh size h for the Crouzeix–Raviart pair, which in a confined plastic zone has no spurious mode among its twenty softest where the constant-pressure elements retain between five and nineteen. On Cook’s membrane in three dimensions at Poisson’s ratio 0.4999, on three meshes of 327 to 14 475 elements, TET15-P1 and the composite tetrahedron agree in the tip displacement within 0.6% on every mesh, and their pressure fields are smooth away from the singular clamped edge, where the pressure of the quadratic tetrahedron without projection changes sign inside elements over the whole membrane; in the plastic case at Poisson’s ratio 0.29 every element converges to the same tip displacement within 0.5% on the finest mesh. A fourteen-point quadrature rule exact to degree five is the smallest that leaves the element its six rigid-body modes, the row-sum lumped mass of the nodal basis is positive, and the explicit stable step is 3.3 times smaller than that of the linear tetrahedron. On the Taylor bar impact in explicit dynamics, with plastic strains above 3, TET15-P1 gives the final radius within 0.05% of the converged value on 25 029 elements, where the composite tetrahedron needs 1.5 million elements for 0.09%, at 1/67 of its core-seconds on the same machine; per element and time step on one core it takes 1.40 times as long. The smoothness of the equivalent plastic strain, the accumulated norm of the plastic strain increments, in problems with localized plastic flow is not measured here.

1 Problem statement

Tetrahedra are the only practical choice for meshing complex geometry automatically. Their accuracy is insufficient in four regimes: near-incompressible response at bulk-to-shear modulus ratios $\kappa/\mu \lesssim 10^4$, with κ the bulk modulus and μ the shear modulus; isochoric (volume-preserving) plasticity; large rotation; and weak shocks, pressure waves of small amplitude, in materials of high κ/μ . The target is one spatial formulation for all four, with the time integrator the only difference between the quasi-static and the dynamic regimes.

Two established remedies fail, and each failure is the failure of one hypothesis of the same theorem; this section identifies the hypothesis in each case.

Constant pressure relieves locking and introduces soft modes. The composite tetrahedron, a ten-node tetrahedron subdivided into twelve linear tetrahedra with one element-constant pressure [21, 28], and its five-field extension [14] relieve volumetric locking — the stiffening of a discretization whose displacement space contains too few discrete isochoric fields — with that pressure. Foulk III et al. [14] report that this choice — and it is the choice, not the element — produces a deleterious soft mode “common to both the quadratic and composite tetrahedral element with constant pressure formulations,” remedied there by a convex energy penalty. The mechanism is a coercivity failure on the discretely isochoric subspace: an element-constant multiplier does not resist displacement patterns whose volume change varies within the element and averages to zero, so those patterns carry only shear energy. In elasticity every projected-pressure element has such patterns, and a stable pair keeps their energy bounded away from zero; in a confined plastic zone the shear response softens along the flow direction and the pattern can approach zero energy. §9.4 shows that on the quadratic tetrahedron these patterns

are not exactly of zero energy, and §6 states the question that count does not settle.

Residual-based stabilization contaminates the deviator. Variational-multiscale stabilization of linear tetrahedra [23] introduces a subgrid pressure of the form

$$p' = c_1 \tau_s (\kappa \operatorname{Div} \dot{\boldsymbol{\varphi}} - \dot{p}), \quad \tau_s = c_2 \Delta t, \quad (1)$$

with c_1 and c_2 tuned coefficients, which is added to the stress before the constitutive update is evaluated. Two consequences follow. The spatial solution depends on the time step Δt through τ_s , so there is no quasi-static limit independent of it. The coefficient c_2 scales a term that reaches the deviator, and therefore the plastic return map — the algorithm that projects a trial stress onto the yield surface — which is the origin of the oscillations in the equivalent plastic strain, the accumulated norm of the plastic strain increments. No choice of c_2 removes this, because the term enters the deviator by construction.

The common cause. Brezzi’s theorem for the mixed problem asks two things of the pair — the discrete space U_h for the displacement together with the discrete space Q_h for the pressure [6, §3.4, §5.2]: the inf-sup condition, that every pressure mode be resisted by some displacement, whose failure is locking on one side and spurious pressure on the other; and coercivity of the deviatoric form on the discretely isochoric subspace $\ker \boldsymbol{G}_h$, that no displacement the pressure does not resist have zero deviatoric energy, whose failure is a soft mode. The first is a property of the pair alone and is settled in §9.5. The second depends on the material tangent as well as on the pair, and §6 states what this note claims about it. Neither difficulty is a stabilization problem: the remedy for the first is a pair satisfying the inf-sup condition uniformly in the mesh size, and the second must then be checked where the material can violate it. §2 defines both hypotheses.

Organization. §2 defines volumetric locking, the mixed problem, the inf-sup condition and Brezzi’s theorem, and gives the construction by which bubble functions satisfy the inf-sup condition. §3 states the continuum formulation and its reduction to mean dilatation. §4 defines the discrete spaces, the element, its quadrature rule and mass, and the element algorithm in the form an implementation needs. §5 gives the plasticity model. §6 states the claim about soft modes under a plastic tangent, §7 the explicit time integration, and §8 an optional extension. §9 reports the measurements on the linearized operator, §10 the nonlinear results on Cook’s membrane, §11 the Taylor bar impact in explicit dynamics, and §12 and §13 summarize the evidence, the open measurements and the limitations.

2 Background: locking, the mixed problem and the inf-sup condition

This section defines the terms the rest of the note uses. The setting is small-strain isotropic elasticity, because every stability property below is measured on the linearized operator at the undeformed state (§4.1); the notation $\boldsymbol{\epsilon}(\boldsymbol{v}) := \operatorname{sym} \operatorname{Grad} \boldsymbol{v}$ is the small strain of a displacement $\boldsymbol{v}$, $\operatorname{dev} \boldsymbol{\epsilon} := \boldsymbol{\epsilon} - \frac{1}{3}(\operatorname{tr} \boldsymbol{\epsilon})\boldsymbol{I}$ its deviatoric part, and $\operatorname{Div} \boldsymbol{v} = \operatorname{tr} \boldsymbol{\epsilon}(\boldsymbol{v})$ its volume change.

2.1 Why low-order tetrahedra lock

The stored energy of an isotropic linear elastic material,

$$W(\boldsymbol{\epsilon}) = \mu |\operatorname{dev} \boldsymbol{\epsilon}|^2 + \frac{\kappa}{2} (\operatorname{tr} \boldsymbol{\epsilon})^2, \quad (2)$$

assigns the energy $\kappa/2$ per unit squared volume change and μ per unit squared shape change. As κ/μ grows, minimizing the energy over a discrete displacement space drives $\operatorname{Div} \mathbf{u}_h$ toward zero at every quadrature point, a point of the integration rule: in the limit the discrete solution is confined to the discrete displacements whose volume change vanishes at every quadrature point. Whether that subspace is rich enough to represent the deformation is a matter of counting. The divergence of a linear tetrahedron is constant, so each element imposes one constraint; a mesh with six tetrahedra per vertex, the count of the structured meshes of §9.1, has about six constraints per vertex against three displacement unknowns per vertex, two constraints per unknown. Few displacements satisfy all of them, so the discrete solid is far stiffer than the material, by an amount that grows with κ/μ and does not decrease with the mesh size. That is volumetric locking. The quadratic tetrahedron has a linear divergence, four constraints per element against about $2 \cdot 3 = 6$ unknowns per element on the same meshes (§9.3), and locks as well.

The remedy is to impose the constraint in an average sense, through an independent pressure field that acts as its Lagrange multiplier, the force that enforces it. This is a mixed method: a method with two independent discrete fields, displacement and pressure, whose behavior depends on the pair of discrete spaces chosen for them.

2.2 The mixed problem and the two hypotheses of Brezzi's theorem

In the incompressible limit the mixed problem reads: find the displacement $\mathbf{u}_h \in U_h$ and the pressure $p_h \in Q_h$ with

$$\begin{aligned} a(\mathbf{u}_h, \mathbf{v}) + b(\mathbf{v}, p_h) &= f(\mathbf{v}) \quad \text{for all } \mathbf{v} \in U_h, \\ b(\mathbf{u}_h, q) - \frac{1}{\kappa} (p_h, q) &= 0 \quad \text{for all } q \in Q_h, \end{aligned} \quad (3)$$

with

$$a(\mathbf{u}, \mathbf{v}) := 2\mu \int_{\Omega} \operatorname{dev} \boldsymbol{\epsilon}(\mathbf{u}) : \operatorname{dev} \boldsymbol{\epsilon}(\mathbf{v}) \, dV, \quad b(\mathbf{v}, q) := \int_{\Omega} q \operatorname{Div} \mathbf{v} \, dV, \quad (4)$$

with f the work of the external loads and $(\cdot, \cdot)$ the L^2 inner product; a is the deviatoric energy form, b the work of the pressure on the volume change, and the second equation is the constitutive relation $p = \kappa \operatorname{Div} \mathbf{u}$ tested against the pressure space, which becomes the constraint $b(\mathbf{u}_h, q) = 0$ as $\kappa \rightarrow \infty$. Brezzi's theorem [7], [6, Theorem 5.2.1 and §5.2], gives two conditions under which (3) has a unique solution that converges to the continuum one with constants independent of the mesh size h .

The inf-sup condition. There is a $\beta > 0$, independent of h , with

$$\beta_h := \inf_{q \in Q_h} \sup_{\mathbf{v} \in U_h} \frac{\int_{\Omega} q \operatorname{Div} \mathbf{v} \, dV}{\|q\|_{L^2} \|\operatorname{Grad} \mathbf{v}\|_{L^2}} \geq \beta. \quad (5)$$

In words: every discrete pressure does work on the volume change of some discrete displacement, in a ratio bounded below independently of the mesh. The condition is also named after Ladyzhenskaya, Babuška and Brezzi (LBB). Its failure has two forms. If $\beta_h = 0$ there is a pressure mode that no displacement resists: the pressure is not determined, modes of alternating sign between neighboring elements appear, and the system becomes singular as $\kappa/\mu \rightarrow \infty$; the pair P_1/P_0 (linear displacement, constant pressure) has 45 such modes on the coarsest mesh of §9.5 and 558 on the finest. If $\beta_h > 0$ on every mesh but $\beta_h \rightarrow 0$ under refinement, the pressure is determined but the error grows with $1/\beta_h$: the stresses oscillate and the displacement is too stiff, and locking returns as the mesh is refined.

Coercivity on the discretely isochoric subspace. There is an $\alpha > 0$, independent of h , with

$$a(\mathbf{v}, \mathbf{v}) \geq \alpha \|\text{Grad } \mathbf{v}\|_{L^2}^2 \quad \text{for all } \mathbf{v} \in \ker \mathbf{G}_h := \{\mathbf{v} \in U_h : b(\mathbf{v}, q) = 0 \text{ for all } q \in Q_h\}. \quad (6)$$

In words: every nonzero discrete displacement whose volume change the pressure space does not see has a deviatoric energy bounded below. A displacement in $\ker \mathbf{G}_h$ with zero deviatoric energy has no stiffness at all; it is a spurious soft mode.

The error estimate. Under both hypotheses [6, Theorem 5.2.2],

$$\|\text{Grad}(\mathbf{u} - \mathbf{u}_h)\| + \|p - p_h\| \leq C \left(1 + \frac{1}{\alpha}\right) \left(1 + \frac{1}{\beta_h}\right) \left[\inf_{\mathbf{v} \in U_h} \|\text{Grad}(\mathbf{u} - \mathbf{v})\| + \inf_{q \in Q_h} \|p - q\| \right] : (7)$$

the error is at most a factor times the best approximation the two spaces allow, and the factor is what the two constants control.

Measuring β_h . With $\{\phi_i\}$ a basis of U_h , the functions of nodes on the fixed boundary excluded, and $\{\chi_m\}$ a basis of Q_h , let

$$K_{ij} = \int_{\Omega} \text{Grad } \phi_i : \text{Grad } \phi_j \, dV, \quad G_{mi} = \int_{\Omega} \chi_m \text{Div } \phi_i \, dV, \quad M_{mn} = \int_{\Omega} \chi_m \chi_n \, dV, \quad (8)$$

the Gram matrix of the norm $\|\text{Grad } \mathbf{v}\|$, the discrete constraint operator and the pressure mass matrix. For a fixed pressure vector q the largest value over $\mathbf{v}$ of $(q^\top \mathbf{G} \mathbf{v})^2 / \mathbf{v}^\top \mathbf{K} \mathbf{v}$ is $q^\top \mathbf{G} \mathbf{K}^{-1} \mathbf{G}^\top q$, so

$$\beta_h^2 = \lambda_{\min}^{\neq 0}(\mathbf{G} \mathbf{K}^{-1} \mathbf{G}^\top, \mathbf{M}), \quad (9)$$

the smallest nonzero generalized eigenvalue [9]. With the whole boundary fixed the constant pressure does no work on any admissible displacement, so one eigenvalue is always zero and is excluded; any further zero is an unresisted pressure mode. §9.1 gives the computation.

2.3 How the enriched pair satisfies the inf-sup condition

A sufficient condition for (5) is the existence of a Fortin operator [13], [6, §5.4]: a linear map Π_h from the continuum displacement space into U_h with

$$b(\mathbf{v} - \Pi_h \mathbf{v}, q) = 0 \text{ for all } q \in Q_h, \quad \|\text{Grad } \Pi_h \mathbf{v}\| \leq C_\Pi \|\text{Grad } \mathbf{v}\|, \quad (10)$$

which gives $\beta_h \geq \beta/C_\Pi$ with β the continuum inf-sup constant. The operator must reproduce, for every q in Q_h , the work of q on the volume change of $\mathbf{v}$, with a bounded gradient. For a pressure that is linear on each element T , a tetrahedron of the mesh, and discontinuous between elements, integration by parts on one element,

$$\int_T q \operatorname{Div} \mathbf{w} \, dV = - \int_T \operatorname{Grad} q \cdot \mathbf{w} \, dV + \int_{\partial T} q \mathbf{w} \cdot \mathbf{n} \, dS, \quad (11)$$

shows what has to be matched: the flux $\mathbf{w} \cdot \mathbf{n}$ weighted by q on each face, and the interior integral of $\mathbf{w}$ against the constant $\operatorname{Grad} q$. Two kinds of bubble supply the two parts. A face bubble b_F , a cubic polynomial that vanishes on the three other faces of the element, times the face normal $\mathbf{n}_F$, changes the boundary term on face F alone and controls the mean of q on that face, and with it the jump of the pressure between the two elements sharing F ; four face bubbles per element control the element mean of q and its values on the faces. The interior bubble b_T , a quartic polynomial that vanishes on the whole element boundary, times $\operatorname{Grad} q$, changes the interior term alone: $\int_T q \operatorname{Div}(b_T \operatorname{Grad} q) \, dV = - \int_T b_T |\operatorname{Grad} q|^2 \, dV < 0$ for $\operatorname{Grad} q \neq 0$, so it controls the gradient of the pressure inside each element. Face bubbles alone leave the linear part of the pressure uncontrolled, and the interior bubble alone leaves the constant part uncontrolled: it vanishes on the boundary, so its volume change has zero mean on the element and does no work on a constant pressure (Remark 7). The displacement space with both enrichments over the linear discontinuous pressure is the three-dimensional Crouzeix–Raviart pair [11], [6, Example 8.7.2], whose stability in three dimensions, for the whole family of Crouzeix–Raviart elements, is also proved by Sauter and Torres [22], and the construction above is the proof of its stability on any shape-regular family of meshes. §9.5 measures its constant and confirms that neither enrichment alone suffices.

3 Continuum formulation

This section states the kinematics, the stored energy, the three-field functional, its reduction to a displacement functional, the requirements it places on the material, and the stress it produces.

3.1 Kinematics and the volumetric–isochoric split

A body $\Omega \subset \mathbb{R}^3$ undergoes a motion $\mathbf{x} = \boldsymbol{\varphi}(\mathbf{X}, t)$ with deformation gradient $\mathbf{F} := \operatorname{Grad} \boldsymbol{\varphi}$ and Jacobian $J := \det \mathbf{F} > 0$. The multiplicative factorization

$$\mathbf{F} = J^{1/3} \bar{\mathbf{F}}, \quad \bar{\mathbf{F}} := J^{-1/3} \mathbf{F}, \quad \det \bar{\mathbf{F}} = 1, \quad (12)$$

separates the two kinematic contributions this formulation treats differently: the mixed treatment acts on J alone, and the plasticity on $\bar{\mathbf{F}}$ alone.

The mixed field is a scalar function of J alone. Let

$$\theta = \theta(J), \quad \theta'(J) > 0, \quad \theta(1) = 0, \quad \theta'(1) = 1, \quad (13)$$

be a strictly increasing volumetric strain measure, normalized to vanish with unit slope in the undeformed state. The normalization has two consequences. It makes κ in (15) the physical bulk modulus for every admissible θ , so that two instances of the formulation are calibrated to the same material rather than to two different ones. And since $\partial J / \partial \mathbf{F} = J \mathbf{F}^{-\top}$,

$$\frac{\partial \theta}{\partial \mathbf{F}} = \theta'(J) J \mathbf{F}^{-\top}, \quad (14)$$

so (13) also forces $\partial\theta/\partial\mathbf{F} = \mathbf{I}$ at $\mathbf{F} = \mathbf{I}$ for every admissible θ (§4.1).

Two choices are used below and are developed in §3.7: the logarithmic measure $\theta = \log J$, and $\theta = J - 1$. Both satisfy (13), and by Remark 1 they nevertheless give different elements.

3.2 Stored energy

The formulation is developed for a stored energy that splits exactly along (12), with a volumetric part quadratic in θ :

$$W(\mathbf{F}) = W^{\text{vol}}(\theta) + W^{\text{iso}}(\bar{\mathbf{F}}), \quad W^{\text{vol}}(\theta) = \frac{\kappa}{2}\theta^2. \quad (15)$$

§3.5 states what each half of (15) requires of the material.

Given (15), the volumetric constitutive relation is linear,

$$p = \partial W^{\text{vol}}/\partial\theta = \kappa\theta, \quad (16)$$

and it is this linearity — not the choice of θ , and not the logarithmic measure specifically — that makes the discrete pressure equal $\kappa\bar{\theta}$ pointwise in §4.2. The elimination itself rests on less, as §3.5 shows.

The isochoric response is the W^{iso} the material supplies. It is evaluated at the pointwise $\bar{\mathbf{F}}$ and is never projected, which is the property §5 requires. Plasticity acts entirely inside W^{iso} , through an isochoric internal variable evolved by a return map; neither (15) nor (16) changes when it does.

3.3 Three-field functional

Only the volumetric response is treated mixedly. Introduce an independent volumetric strain $\bar{\theta}$ and constrain it to $\theta(\boldsymbol{\varphi})$ through a multiplier $\bar{p}$, which at stationarity is the pressure:

$$\begin{aligned} \Pi[\boldsymbol{\varphi}, \bar{\theta}, \bar{p}] := & \int_{\Omega} W^{\text{vol}}(\bar{\theta}) \, dV + \int_{\Omega} W^{\text{iso}}(\bar{\mathbf{F}}(\boldsymbol{\varphi})) \, dV + \int_{\Omega} \bar{p} (\theta(\boldsymbol{\varphi}) - \bar{\theta}) \, dV \\ & - \int_{\Omega} \rho_0 \mathbf{b} \cdot \boldsymbol{\varphi} \, dV - \int_{\partial_T \Omega} \mathbf{T} \cdot \boldsymbol{\varphi} \, dS, \end{aligned} \quad (17)$$

with ρ_0 the reference density, $\mathbf{b}$ the body force per unit mass, and $\mathbf{T}$ the traction prescribed on the part $\partial_T \Omega$ of the boundary. The isochoric energy is evaluated at the pointwise $\bar{\mathbf{F}}(\boldsymbol{\varphi})$ — for the Hencky instance of §3.7, at $\mathbf{e}(\boldsymbol{\varphi})$ — and carries no auxiliary field; Remark 3 gives the reason.

3.4 Reduction to the mean-dilatation formulation

With $\bar{\theta}$ and $\bar{p}$ in one discrete space, (17) reduces to the mean-dilatation formulation of Simo et al. [26]; this section gives the reduction.

Put θ and $\bar{p}$ in the same discrete space Q_h , as §4.1 requires. Stationarity in $\bar{\theta}$ then gives $\bar{p} = \kappa\bar{\theta}$ exactly rather than after a projection, and stationarity in $\bar{p}$ gives $\bar{\theta} = \mathcal{P}_h \theta$, the L^2 projection of θ onto Q_h . The constraint term of (17) then vanishes identically:

$$\int_{\Omega} \bar{p} (\theta - \bar{\theta}) \, dV = \kappa \int_{\Omega} \bar{\theta} (\theta - \mathcal{P}_h \theta) \, dV = 0, \quad (18)$$

because $\kappa\bar{\theta} \in Q_h$ while $(\theta - \mathcal{P}_h\theta) \perp Q_h$ by the definition of the projection. What is left is a functional in the motion alone,

$$\Pi[\varphi] = \int_{\Omega} W^{\text{iso}}(\bar{\mathbf{F}}(\varphi)) \, dV + \int_{\Omega} \frac{\kappa}{2} (\mathcal{P}_h\theta)^2 \, dV, \quad (19)$$

whose stationarity conditions are exactly the three-field equations of §3.8.

Equation (19) is the mean-dilatation formulation of Simo et al. [26], in which the volumetric strain is replaced by its projection onto an element-wise space, and the mechanism of the $\bar{B}$ family (projected strain–displacement operator) and the $\bar{F}$ family (projected deformation gradient) of elements; Boffi et al. [6, §8.12] set out the relation between reduced integration, projected pressure, and the mixed method it is equivalent to. The reduction is exact by (18), and no property of θ or of W^{vol} enters it: the exactness is a property of the discretization, that the two auxiliary fields share one space. What the material must supply is stated in §3.5, and the quantity the formulation leaves open is Q_h , which §4.1 selects and §9 measures.

3.5 Requirements on the material

This section states what each half of (15) requires of the material, together with a third condition that is not a requirement of this section, and reports which materials meet them. The two conditions are those of the form of the element implemented, in which the material is called once per quadrature point and the pressure equals $\kappa\bar{\theta}$ at every point; §3.6 gives the form of the element for a material that satisfies neither, at the price of one projection and one pass more.

An exact volumetric–isochoric split. This condition is the one that lets the element call the material at the pointwise deformation. The element evaluates W at the pointwise deformation, subtracts the pointwise volumetric energy and adds it back at the projected strain, so the subtraction must leave no volumetric remainder:

$$W(\mathbf{F}) = W^{\text{vol}}(J) + W^{\text{iso}}(\bar{\mathbf{F}}), \quad \bar{\mathbf{F}} := J^{-1/3}\mathbf{F}, \quad \det \bar{\mathbf{F}} = 1. \quad (20)$$

Whether a material satisfies (20) is decidable without knowing W^{vol} : if the split exists then $W(\mathbf{F}) - W(\bar{\mathbf{F}}) = W^{\text{vol}}(J) - W^{\text{vol}}(1)$ depends on J alone, so holding J fixed while varying the isochoric part must leave the difference unmoved. Column 2 of Table 1 is that test, run over seven material models: linear elasticity, Saint-Venant–Kirchhoff, neo-Hookean, reciprocal neo-Hookean, Seth–Hill, Hencky and the J_2 plasticity model of Simo and Hughes [25], as implemented in the finite-element code Norma. Five of the seven satisfy it.

A quadratic W^{vol} . This condition is a convenience and not a necessity. Stationarity in $\bar{p}$ gives $\bar{\theta} = \mathcal{P}_h\theta$ regardless of the material, and stationarity in $\bar{\theta}$ then gives

$$\bar{p} = \mathcal{P}_h \left(\partial W^{\text{vol}} / \partial \bar{\theta} \right) \quad (21)$$

in general — the projection of a known function of the already computed $\bar{\theta}$, one more mass-matrix solve with a nonlinear integrand and nothing implicit. When W^{vol} is quadratic its derivative is linear, $\kappa\bar{\theta}$ already lies in Q_h , the projection in (21) is the identity, and $\bar{p} = \kappa\bar{\theta}$ holds pointwise; otherwise $\bar{p} = \mathcal{P}_h W^{\text{vol}'}(\mathcal{P}_h\theta)$ is a doubly projected pressure. The reduction (19) is exact either way, with $W^{\text{vol}}(\mathcal{P}_h\theta)$ in place of the quadratic, and the elimination of §4.2 is linear either way; what the quadratic case gives is the pointwise relation and one projection fewer.

Table 1: What seven material models satisfy. Column 2 is the spread of $W(\mathbf{F}) - W(\bar{\mathbf{F}})$ across twenty random isochoric parts at each of four fixed $J \in [0.85, 1.2]$, relative to its own magnitude; a spread below 10^{-12} means the split (20) holds to rounding error. Columns 3 and 4 are the relative residual of a least-squares fit of the extracted $W^{\text{vol}}(J) = W(J^{1/3}\mathbf{I}) - W(\mathbf{I})$ over $J \in [0.8, 1.25]$ to a quadratic in each candidate volumetric variable; in both exact cases the fitted coefficient is $\kappa/2$ to five digits. Two materials admit a quadratic W^{vol} , and they disagree on the variable. Script: `prototype/materials.jl`, which needs Norms.

material	split (20)	W^{vol} quadratic in	
	spread at fixed J	$J - 1$	$\log J$
Linear elastic	2.0×10^0	—	—
Saint-Venant Kirchhoff	3.0×10^1	—	—
Neo-Hookean	4.2×10^{-13}	5.9×10^{-2}	1.3×10^{-1}
Reciprocal neo-Hookean	4.8×10^{-13}	1.8×10^{-1}	6.5×10^{-3}
Seth–Hill ($m = n = 1$)	4.3×10^{-13}	1.8×10^{-1}	6.5×10^{-3}
Hencky	1.4×10^{-13}	1.8×10^{-1}	1.7×10^{-15}
J_2 , Simo–Hughes	5.9×10^{-13}	5.3×10^{-15}	1.8×10^{-1}

The logarithmic measure itself. This condition is required by nothing in the present section. It is required by §5.

The case developed here. The formulation is developed in detail for materials satisfying (20) with a quadratic W^{vol} , because that is where the pointwise relation holds. Table 1 admits exactly two, and they do not agree on the variable: Hencky is quadratic in $\log J$ to 1.7×10^{-15} , and the Simo–Hughes J_2 model, the one implemented in Carina, is quadratic in $J - 1$ to 5.3×10^{-15} . Each is exact in its own variable and wrong by 18% in the other, so material and volumetric variable are not independently selectable, and neither choice generalizes to the other material. §3.7 develops both.

The three remaining split materials — neo-Hookean, its reciprocal, and Seth–Hill — are admissible with one projection more (§3.6) and are not developed further.

3.6 General materials

This section gives the element for a material that does not satisfy the conditions of §3.5. Two cases are distinguished by what the material exposes. In the first the split (20) exists and is accessible, and only the volumetric energy $W^{\text{vol}}(\theta)$ is not quadratic. In the second nothing is accessible beyond what every finite-element code obtains from a material model: the stored energy $W(\mathbf{F})$, the first Piola–Kirchhoff stress $\mathbf{P}(\mathbf{F}) = \partial W / \partial \mathbf{F}$ and the tangent $\mathbb{A}(\mathbf{F}) = \partial^2 W / \partial \mathbf{F} \partial \mathbf{F}$, with the internal variables the model carries. The second case is the general one: it applies to a code whose material models cannot be modified, and it contains the first and the split form as special cases.

Case 1: an accessible split with a non-quadratic volumetric energy. Nothing in §3.3 changes. With $\bar{\theta}$ and $\bar{p}$ in one space Q_h , stationarity in $\bar{p}$ gives $\bar{\theta} = \mathcal{P}_h \theta$ as before, and

stationarity in $\bar{\theta}$ gives

$$\int_{\Omega} (W^{\text{vol}'}(\bar{\theta}) - \bar{p}) q \, dV = 0 \text{ for all } q \in Q_h, \quad \text{that is} \quad \bar{p} = \mathcal{P}_h W^{\text{vol}'}(\bar{\theta}), \quad (22)$$

the projection of a known function of the already computed $\bar{\theta}$. The constraint term of (17) vanishes as in (18), because $\bar{p}$ lies in Q_h and $\theta - \mathcal{P}_h \theta$ is orthogonal to it, so the functional reduces to

$$\Pi[\varphi] = \int_{\Omega} W^{\text{iso}}(\bar{\mathbf{F}}(\varphi)) \, dV + \int_{\Omega} W^{\text{vol}}(\mathcal{P}_h \theta) \, dV, \quad (23)$$

whose derivative reproduces (22), since $\mathcal{P}_h$ is self-adjoint in L^2 : $\int W^{\text{vol}'}(\bar{\theta}) \mathcal{P}_h(\delta\theta) \, dV = \int \mathcal{P}_h W^{\text{vol}'}(\bar{\theta}) \delta\theta \, dV$. The stress at a quadrature point is (43) with this $\bar{p}$. Four things follow.

1. With the constant projection nothing changes: $\bar{\theta}$ is constant on the element, so $W^{\text{vol}'}(\bar{\theta})$ is constant and the projection in (22) is the identity. This is the mean-dilatation element of Simo et al. [26], which was formulated for a general W^{vol} .
2. With the linear projection the pressure is projected twice, and $\bar{p}(\xi_q) \neq W^{\text{vol}'}(\bar{\theta}(\xi_q))$: the pressure at a quadrature point is not the constitutive response of the projected strain at that point. Using $W^{\text{vol}'}(\bar{\theta}(\xi_q))$ there instead is not the derivative of any functional, gives an unsymmetric tangent and fails the consistency checks of §4.5; it is not admissible.
3. The tangent replaces the constant matrix $\kappa \mathbf{D} \mathbf{M}^{-1} \mathbf{D}^{\top}$ of §4.4 by $\mathbf{D} \mathbf{M}^{-1} \tilde{\mathbf{U}} \mathbf{M}^{-1} \mathbf{D}^{\top}$ with $\tilde{U}_{mn} = \sum_q w_q j_q W^{\text{vol}''}(\bar{\theta}_q) \chi_m(\xi_q) \chi_n(\xi_q)$, a weighted mass matrix assembled at every iteration, and the volumetric stress $\bar{p}_q \partial \mathbf{g} / \partial \mathbf{F}$ of (44) uses the projected $\bar{p}_q$; the tangent stays symmetric. In the linearized problem of §2.2 the compressibility term $-(1/\kappa)(p, q)$ becomes a term weighted by $1/W^{\text{vol}''}(\bar{\theta})$, which varies inside the element and with the deformation; the inf-sup constant β_h is unchanged, since it depends only on the spaces and on the constraint operator, and the constants of Brezzi's theorem then involve the smallest and largest values of $W^{\text{vol}''}$ over the body, which must be positive: W^{vol} must be convex in θ over the range of $\bar{\theta}$ reached.
4. The second projection can be avoided by a change of variable. Any $W^{\text{vol}}(J)$ with a single minimum at $J = 1$ and $W^{\text{vol}''}(1) = \kappa > 0$ is quadratic in

$$\theta(J) := \text{sign}(J - 1) \sqrt{2W^{\text{vol}}(J)/\kappa}, \quad (24)$$

which satisfies (13) with $\theta = J - 1 + O((J - 1)^2)$ and is convex-admissible wherever W^{vol} is. The element of §4.4 then applies unchanged with this θ . For the volumetric law $W^{\text{vol}} = \frac{\kappa}{4}(J^2 - 1 - 2 \log J)$ of Albany-LCM, for example, $\theta^2 = (J^2 - 1 - 2 \log J)/2$. What distinguishes this choice from a non-quadratic W^{vol} in the variable $J - 1$ is the quantity averaged over the element, $\mathcal{P}_h \theta$ against $\theta(\mathcal{P}_h(J - 1))$; by Remark 1 the two are different elements that converge to the same solution.

Case 2: a material without an accessible split. The element calls the material at a modified deformation gradient whose volume is the projected one. With J the pointwise Jacobian and

$$\tilde{J} := \theta^{-1}(\mathcal{P}_h \theta(J)), \quad s := (\tilde{J}/J)^{1/3}, \quad \tilde{\mathbf{F}} := s \mathbf{F}, \quad \det \tilde{\mathbf{F}} = \tilde{J}, \quad (25)$$

the element is the stationarity of

$$\Pi[\varphi] = \int_{\Omega} W(\tilde{\mathbf{F}}(\varphi)) \, dV - \Pi_{\text{ext}}. \quad (26)$$

Here θ^{-1} is the inverse of the strictly increasing $\theta(J)$; for $\theta = J - 1$, $\tilde{J} = 1 + \mathcal{P}_h(J - 1) = \mathcal{P}_h J$, and for $\theta = \log J$, $\tilde{J} = \exp(\mathcal{P}_h \log J)$. The isochoric part of $\tilde{\mathbf{F}}$ is that of $\mathbf{F}$, since $\tilde{J}^{-1/3} \tilde{\mathbf{F}} = J^{-1/3} \mathbf{F}$, so the material is evaluated at the pointwise shape change and the projected volume change. Equation (26) is a functional of the motion alone for every material, whether or not its energy splits; it is the mean-dilatation construction of Simo et al. [26] with the projection $\mathcal{P}_h$ in place of the element mean, and with the constant projection and $\theta = J - 1$ it is that construction, in which the $\bar{\mathbf{F}}$ element of de Souza Neto et al. [12] replaces the element mean of J by its value at the centroid. For a material with the split (20), $W(\tilde{\mathbf{F}}) = W^{\text{vol}}(\theta(\tilde{J})) + W^{\text{iso}}(\tilde{\mathbf{F}}) = W^{\text{vol}}(\mathcal{P}_h \theta) + W^{\text{iso}}(\tilde{\mathbf{F}})$, so (26) is (23), and with a quadratic W^{vol} it is (19): the general element contains both.

The residual of case 2. Let $\tilde{\mathbf{P}} := \mathbf{P}(\tilde{\mathbf{F}})$ be the stress the material returns at $\tilde{\mathbf{F}}$, with its internal variables updated there, and let

$$\tilde{p} := \frac{1}{3\tilde{J}} \tilde{\mathbf{P}} : \tilde{\mathbf{F}} = \frac{1}{3} \text{tr } \tilde{\boldsymbol{\sigma}} \quad (27)$$

be the mean Cauchy stress at $\tilde{\mathbf{F}}$, positive in expansion. The variation of (25) is $\delta \tilde{\mathbf{F}} = s \delta \mathbf{F} + \delta s \mathbf{F}$ with $\delta s = \frac{s}{3} (\delta \tilde{J}/\tilde{J} - \delta J/J)$, $\delta J = \mathbf{J} \mathbf{F}^{-\top} : \delta \mathbf{F}$ and $\delta \tilde{J} = \delta \bar{\theta}/\theta'(\tilde{J})$, $\delta \bar{\theta} = \mathcal{P}_h(\theta'(J) \delta J)$. Substituting into $\delta \Pi = \int \tilde{\mathbf{P}} : \delta \tilde{\mathbf{F}} \, dV$, using $\tilde{\mathbf{P}} : \tilde{\mathbf{F}} = 3\tilde{J}\tilde{p}$ and moving the projection onto the other factor by self-adjointness gives the stress that multiplies $\text{Grad } \boldsymbol{\psi}$ at a quadrature point,

$$\mathbf{P}_q = s_q \tilde{\mathbf{P}}_q - \tilde{p}_q \tilde{J}_q \mathbf{F}_q^{-\top} + \bar{p}_q \theta'(J_q) J_q \mathbf{F}_q^{-\top}, \quad \bar{p} := \mathcal{P}_h \left(\frac{\tilde{p}}{\theta'(\tilde{J})} \right), \quad (28)$$

and $\mathbf{R}_a = \sum_q w_q j_q \mathbf{P}_q \text{Grad } N_a(\boldsymbol{\xi}_q)$ as in (43). The first term is the material's stress scaled to the pointwise volume, the second removes its mean part, and the third puts back the projected mean stress with the pointwise volumetric kinematics. For a material with the split, $\tilde{\mathbf{P}} = \mathbf{P}^{\text{iso}}(\tilde{\mathbf{F}}) s^{-1} + W^{\text{vol}'}(\bar{\theta}) \theta'(\tilde{J}) \tilde{J} \tilde{\mathbf{F}}^{-\top}$, the isochoric stress does no work on $\tilde{\mathbf{F}}$, and (28) reduces to (43) with $\bar{p} = \mathcal{P}_h W^{\text{vol}'}(\bar{\theta})$; with a quadratic W^{vol} the projection is the identity and $\bar{p} = \kappa \bar{\theta}$. For $\theta = J - 1$ and the constant projection, $\bar{p} = \tilde{p}$ on the element and (28) is $\mathbf{P}_q = s_q \tilde{\mathbf{P}}_q + \tilde{p} (J_q - \tilde{J}) \mathbf{F}_q^{-\top}$.

The passes of case 2. Three passes over the quadrature points replace the two of §4.4. The first computes $\bar{\theta}$ by (42) and, from it, $\tilde{J}_q$ and s_q . The second forms $\tilde{\mathbf{F}}_q$, calls the material once at each point, stores $\tilde{\mathbf{P}}_q$ and $\tilde{p}_q$, and accumulates $\bar{b}_m = \sum_q w_q j_q \chi_m(\boldsymbol{\xi}_q) \tilde{p}_q / \theta'(\tilde{J}_q)$, from which $\bar{p} = \mathbf{M}^{-1} \bar{\mathbf{b}}$ with the mass matrix of (42). The third assembles the residual from the stored $\tilde{\mathbf{P}}_q$ and (28). The storage is one stress tensor per quadrature point of the element; the material is called once per point per residual, as in the split form. Where the split is accessible the second projection is still exact but the third pass is unnecessary, since $\bar{p}$ follows from $\bar{\theta}$ alone.

The tangent of case 2. The tangent is the second derivative of (26). Along two increments with gradients $\delta_1 \mathbf{F}$ and $\delta_2 \mathbf{F}$ at the quadrature points,

$$D^2 \Pi[\delta_1 \boldsymbol{\varphi}, \delta_2 \boldsymbol{\varphi}] = \sum_q w_q j_q \left[\delta_1 \tilde{\mathbf{F}}_q : \mathbb{A}_q : \delta_2 \tilde{\mathbf{F}}_q + \tilde{\mathbf{P}}_q : (\delta_1 s_q \delta_2 \mathbf{F}_q + \delta_2 s_q \delta_1 \mathbf{F}_q + \delta_1 \delta_2 s_q \mathbf{F}_q) \right], \quad (29)$$

with $\mathbb{A}_q$ the material tangent at $\tilde{\mathbf{F}}_q$, the variations $\delta \tilde{\mathbf{F}}$ and δs as above, and

$$\begin{aligned} \delta_1 \delta_2 s &= \frac{\delta_1 s}{3} \left(\frac{\delta_2 \tilde{J}}{\tilde{J}} - \frac{\delta_2 J}{J} \right) + \frac{s}{3} \left(\frac{\delta_1 \delta_2 \tilde{J}}{\tilde{J}} - \frac{\delta_1 \tilde{J} \delta_2 \tilde{J}}{\tilde{J}^2} - \frac{\delta_1 \delta_2 J}{J} + \frac{\delta_1 J \delta_2 J}{J^2} \right), \\ \delta_1 \delta_2 J &= J (\mathbf{F}^{-\top} : \delta_1 \mathbf{F}) (\mathbf{F}^{-\top} : \delta_2 \mathbf{F}) - J \operatorname{tr} (\mathbf{F}^{-1} \delta_1 \mathbf{F} \mathbf{F}^{-1} \delta_2 \mathbf{F}), \\ \delta_1 \delta_2 \tilde{J} &= \frac{\mathcal{P}_h(\theta''(J) \delta_1 J \delta_2 J + \theta'(J) \delta_1 \delta_2 J)}{\theta'(\tilde{J})} - \frac{\theta''(\tilde{J})}{\theta'(\tilde{J})} \delta_1 \tilde{J} \delta_2 \tilde{J}. \end{aligned} \quad (30)$$

The projection in $\delta \tilde{J}$ couples the quadrature points of the element, as $\delta \bar{\theta}$ does in (44): in matrix form the tangent is the sum of the quadrature-point blocks and of terms of the form $\mathbf{D} \mathbf{M}^{-1}(\cdot) \mathbf{M}^{-1} \mathbf{D}^\top$ and $\mathbf{D} \mathbf{M}^{-1}(\cdot)$ with the matrix $\mathbf{D}$ of §4.4. The tangent is symmetric, and an implementation is verified by the difference checks of §4.5: the residual against the energy (26) and the tangent against the residual, which detect any omitted term.

What the general element changes. The requirements of §3.5 disappear: the material supplies W , $\mathbf{P}$ and $\mathbb{A}$ and nothing else, and the choice of $\theta(J)$ is free within (13). The return map runs at $\tilde{\mathbf{F}}_q$, and the internal variables are updated there, so strain and stress are evaluated at the same point and the same deformation, as §5 requires. For a material whose isochoric response depends on J , a material without the split, that response is evaluated at $\tilde{J}$ instead of J ; the difference is of the order of $\theta - \mathcal{P}_h \theta$ and vanishes under refinement. The pressure entering the residual is the projection of the material's own mean stress at $\tilde{\mathbf{F}}$, with the pointwise relation between pressure and projected strain recovered only when W^{vol} is quadratic in θ . The inf-sup constant of the pair is unchanged. Coercivity of the linearized operator requires the material's volumetric tangent, $\partial \bar{p} / \partial \tilde{J}$ at fixed shape, to be positive over the states reached. The element needs the element-level assembly of §4.4 and one pass more, and its cost per element is one material call per quadrature point plus the two 4×4 solves.

Implementation. Both forms are implemented in Carina in one physics type with a form parameter, chosen at setup: the split form when the material exposes the split, the general form otherwise, and the general form on request for any material (input key `volumetric form`). The volumetric variable is a property of the element, not of the material (input key `volumetric strain`, with the values $\log J$ and $J - 1$): a material without the split uses $\log J$ unless the key says otherwise; a material with the split uses its own θ , and a key that differs from it is an error, since the split form's pressure $\kappa \bar{\theta}$ is the stationarity condition of the energy $\frac{\kappa}{2} \theta^2$ written in the material's θ . The general residual takes the three passes above with the fourteen stresses $\tilde{\mathbf{P}}_q$ held on the stack. The tangent is not coded from (29) directly: the element residual is written generically in its number type and differentiated in forward mode, with the material's stress and tangent at $\tilde{\mathbf{F}}_q$ supplying the value and the derivative of $\tilde{\mathbf{P}}_q$ along each direction, so that a material that writes its internal variables into real-valued storage is never called with derivative-carrying numbers; this gives (29) with the material's own $\mathbb{A}_q$. The action on a

Table 2: The two instances of (15). Everything above the rule is fixed by the generic development; everything below it is where the instances differ. $\mathbf{b}^e := \mathbf{F}^e(\mathbf{F}^e)^\top$ is the elastic left Cauchy–Green tensor, with $\mathbf{F}^e$ the elastic part of the multiplicative decomposition $\mathbf{F} = \mathbf{F}^e \mathbf{F}^p$, and $\bar{\mathbf{b}}^e := J^{-2/3} \mathbf{b}^e$ its isochoric part.

	Hencky	Simo–Hughes
volumetric variable θ	$\log J$	$J - 1$
$\theta'(J)$	$1/J$	1
$\partial\theta/\partial\mathbf{F}$	$\mathbf{F}^{-\top}$	$J \mathbf{F}^{-\top}$
$W^{\text{vol}}(\theta)$	$\frac{\kappa}{2}(\log J)^2$	$\frac{\kappa}{2}(J - 1)^2$
isochoric argument	$\mathbf{e} = \text{dev } \boldsymbol{\epsilon}, \boldsymbol{\epsilon} = \frac{1}{2} \log \mathbf{C}$	$\bar{\mathbf{b}}^e$
W^{iso}	$\mu \ \mathbf{e}\ ^2$	$\frac{\mu}{2}(\text{tr } \bar{\mathbf{b}}^e - 3)$
return map	radial return in $\boldsymbol{\epsilon}$	radial return on $\bar{\mathbf{b}}^e$
small-strain models usable unmodified	yes	no
implemented in Carina	no	yes, with consistent tangent

vector takes one pass with one derivative direction; the element matrix and the diagonal and nodal-block kernels of the preconditioners take one pass with three directions per node, fifteen passes per element. Measured on the Taylor bar mesh at $h = 0.75$ mm (§11, 339 993 unknowns) with the J_2 model, the general residual takes 1.36 times the split form’s time on a CPU with 12 threads (16.2 against 11.9 ms) and 1.60 times on an AMD RX 7600 (13.6 against 8.5 ms); per explicit step of a run the ratio is 1.36 on a CPU with 16 threads, 1.44 on the RX 7600, 1.60 on an NVIDIA L4, 1.84 on an A100 and 2.05 on a V100 (Table 16). The split form remains the default for that model.

Remark 1 (The volumetric variable is a modeling choice). Projecting $\log J$ and projecting $J - 1$ do not give the same element. Measured on one tetrahedron under a displacement quadratic in $\mathbf{X}$, with the element mean as the projection and a Hencky material so that both variants use the same exact W^{vol} , the element energies differ by 9.4×10^{-4} at 0.5% nominal strain, 1.0×10^{-2} at 5%, 5.6×10^{-2} at 20% and 2.7×10^{-1} at 50% (prototype/materials.jl). The difference is $O(\text{strain})$ and vanishes with it, so both are consistent discretizations of the same continuum problem; they are different elements, and they differ by 5.6% to 27% at 20% to 50% strain, the finite-deformation regime of §1. Which variable is projected is therefore a modeling decision.

3.7 The two admissible instances

Table 1 leaves two materials. They agree on everything §3.1–§3.4 requires; Table 2 lists where they differ.

The Hencky instance is the one (31) defines: with

$$\boldsymbol{\epsilon} := \frac{1}{2} \log \mathbf{C} \in \mathbb{S}, \quad \mathbf{C} := \mathbf{F}^\top \mathbf{F}, \quad (31)$$

evaluated by inverse scaling and squaring with Padé approximants — repeated square roots of $\mathbf{C}$ followed by a rational approximant of the logarithm [3] — the volumetric–deviatoric decomposition

$$\boldsymbol{\epsilon} = \frac{1}{3} \theta \mathbf{I} + \mathbf{e}, \quad \theta = \text{tr } \boldsymbol{\epsilon} = \log J, \quad \mathbf{e} = \text{dev } \boldsymbol{\epsilon}, \quad (32)$$

is orthogonal in the Frobenius inner product, and the split of (15) is inherited from the strain measure rather than imposed on it. The Simo–Hughes instance imposes it instead, through the $J^{-2/3}$ factor in $\bar{\mathbf{b}}^e$; $\det \bar{\mathbf{b}}^e = 1$ then makes W^{iso} independent of J by construction. Both are exact, and Table 1 confirms both to machine precision. Three consequences follow.

The generic development covers both. Nothing in §3.1–§3.4, and nothing in the discretization of §4.1–§4.2, refers to either column of Table 2. The reduction (19), the pointwise relation $\bar{p} = \kappa \bar{\theta}$, and the element-local elimination hold for any θ satisfying (13) and any W^{iso} satisfying (15).

The pairing at $\mathbf{F} = \mathbf{I}$ is instance-independent. By (14), $\partial\theta/\partial\mathbf{F} = \mathbf{I}$ at $\mathbf{F} = \mathbf{I}$ in both columns; §4.1 draws the consequence for the stability results.

The instances differ in plasticity only. The last two rows of Table 2 are the entire difference, and they follow from (32), as §5 develops. Remark 5 states what the implemented instance, Simo–Hughes, entails.

3.8 Variations

With $\varphi \in \mathcal{U}$, the space of motions of finite deformation energy that take the prescribed values on $\partial\varphi\Omega$, the part of the boundary on which the motion is prescribed; $\bar{\theta}, \bar{p} \in \mathcal{Q}$, the space of square-integrable scalar fields; test functions $\psi \in \mathcal{U}$ vanishing on $\partial\varphi\Omega$; and $\vartheta, \pi \in \mathcal{Q}$:

$$\begin{aligned} D_{\varphi}\Pi(\psi) &= \int_{\Omega} \bar{p} \frac{\partial\theta}{\partial\varphi}[\psi] \, dV + \int_{\Omega} \mathbf{s}(e(\varphi)) : \frac{\partial e}{\partial\varphi}[\psi] \, dV - \int_{\Omega} \rho_0 \mathbf{b} \cdot \psi \, dV - \int_{\partial_T\Omega} \mathbf{T} \cdot \psi \, dS = 0, \\ D_{\bar{\theta}}\Pi(\vartheta) &= \int_{\Omega} (\kappa \bar{\theta} - \bar{p}) \vartheta \, dV = 0, \\ D_{\bar{p}}\Pi(\pi) &= \int_{\Omega} (\theta(\varphi) - \bar{\theta}) \pi \, dV = 0. \end{aligned} \tag{33}$$

Here $\mathbf{s}$ is the deviatoric stress conjugate to $\mathbf{e}$, given by the material law (§5). The kinematic derivatives are

$$\frac{\partial\theta}{\partial\varphi}[\psi] = \theta'(J) J \mathbf{F}^{-\top} : \text{Grad } \psi, \quad \frac{\partial e}{\partial\varphi}[\psi] = \mathbb{L}^{\text{dev}}(\mathbf{C}) : \text{sym}(\mathbf{F}^{\top} \text{Grad } \psi), \tag{34}$$

with $\mathbb{L}(\mathbf{C}) := \partial \log \mathbf{C} / \partial \mathbf{C}$ the Fréchet derivative of the matrix logarithm [3] and $\mathbb{L}^{\text{dev}} := \mathbb{L} - \frac{1}{3} \mathbf{I} \otimes \mathbf{C}^{-1}$. The volumetric derivative follows from Jacobi’s formula and (14). The deviatoric expression is written for the Hencky instance of Table 2; the Simo–Hughes instance replaces it with the variation of $\bar{\mathbf{b}}^e$, and nothing outside this equation changes.

3.9 Stress reconstruction

Collecting (33)₁ into the form $\int \mathbf{P} : \text{Grad } \psi \, dV$, with $\mathbf{P}$ the first Piola–Kirchhoff stress, the stress conjugate to $\mathbf{F}$, gives in general

$$\mathbf{P} = \theta'(J) J \bar{p} \mathbf{F}^{-\top} + \mathbf{F} (\mathbb{L}(\mathbf{C}) : \mathbf{s}). \tag{35}$$

The deviatoric term is $\mathbb{L}$ and not $\mathbb{L}^{\text{dev}}$ applied to $\mathbf{s}$. Moving $\mathbb{L}^{\text{dev}}$ off the strain variation in (34) and onto $\mathbf{s}$ takes its transpose, which carries the $\frac{1}{3}$ term onto $\text{tr } \mathbf{s} = 0$, and $\mathbb{L}$ is self-adjoint because $\log \mathbf{C}$ is the gradient of $\text{tr}(\mathbf{C} \log \mathbf{C} - \mathbf{C})$.

When the deviatoric response is elastic, $\mathbf{s} = 2\mu \mathbf{e}$ is a function of the current $\mathbf{C}$ alone through an isotropic law, so $\mathbf{s}$ is coaxial with $\mathbf{C}$ at every quadrature point; the off-diagonal blocks of $\mathbb{L}$ in the principal basis then act on nothing and (35) collapses to

$$\mathbf{P} = (\varpi \mathbf{I} + \mathbf{s}_x) \mathbf{F}^{-\top} = \boldsymbol{\tau} \mathbf{F}^{-\top}, \quad \boldsymbol{\tau} := \varpi \mathbf{I} + \mathbf{s}_x = J \boldsymbol{\sigma}, \quad \varpi := \theta'(J) J \bar{p}, \quad (36)$$

in which $\boldsymbol{\tau} = J \boldsymbol{\sigma}$ is the Kirchhoff stress, $\boldsymbol{\sigma}$ the Cauchy stress, and $\mathbf{s}_x := \mathbf{R} \mathbf{s} \mathbf{R}^\top$ the spatial deviator, with $\mathbf{R}$ the rotation of the polar decomposition $\mathbf{F} = \mathbf{R} \mathbf{U}$. The material $\mathbf{s}$ of (32) and the spatial $\mathbf{s}_x$ are different tensors; in material form the same statement reads $\mathbf{P} = \mathbf{F}^{-\top} \mathbf{s}$, which agrees with (36) because $\mathbf{s}$ and $\mathbf{U}$ commute when coaxial.

The Kirchhoff pressure ϖ carries the whole of the instance dependence. For the Hencky instance $\theta'(J)J = 1$ and $\varpi = \bar{p}$; for the Simo–Hughes instance $\theta'(J)J = J$ and $\varpi = J\bar{p}$, giving $\boldsymbol{\tau} = J\bar{p} \mathbf{I} + \mathbf{s}$ with $\bar{p} = \kappa(J - 1)$, which is the Kirchhoff stress of Simo and Hughes [25, BOX 9.1] and is what both codes compute. Agreement of (35) with both implementations confirms the normalization (13).

Remark 2 (Coaxiality is an elastic property). Equation (36) rests on $\mathbf{s}$ being coaxial with $\mathbf{C}$, which holds only when $\mathbf{s}$ is a function of the current $\mathbf{C}$ alone. It fails in two settings, and the second is the one this formulation exists for. With an independent discrete stress field, $\mathbf{s}$ is tied to $\mathbf{C}$ only weakly and is not coaxial with it at an arbitrary iterate. With plasticity, the return map (46) evaluates $\mathbf{s} = 2\mu(\mathbf{e} - \mathbf{e}_{n+1}^p)$ against a stored plastic strain $\mathbf{e}_n^p$, which is coaxial with the current $\mathbf{C}$ only under proportional loading; whenever the principal axes rotate — the large-rotation regime of §1 — $\mathbf{s}$ is not coaxial with $\mathbf{C}$, the off-diagonal blocks of $\mathbb{L}$ do not vanish, and (36) is wrong. In both settings (36) holds at best at a converged proportional state, and assembling the residual from it gives a residual that is not the derivative of the energy, with an error that is undetected under proportional loading. The residual is assembled from (35); Miehe et al. [18] carry the full projection tensor for this reason.

Remark 3. No deviatoric mixed pair appears because none is needed. Volumetric locking is a constraint phenomenon: at $\kappa/\mu \gg 1$ the near-isochoric constraint over-determines the displacement space, and relieving it requires a multiplier. The deviatoric response is not constrained in the same sense and admits no analogous locking mode. Adding a deviatoric pair adds degrees of freedom, an inf-sup condition that has to be established separately, and — if strain and stress are placed in different spaces — a constitutive law evaluated across two discretizations, which is the error that corrupts plastic strain fields when the strain is interpolated in one space and the stress evaluated in another. §8 gives the construction that adds strain as an independent field should a later requirement demand it.

4 Discretization

This section selects the pair of discrete spaces, defines the element that carries the displacement, gives the element-local elimination of the auxiliary fields, and states the element algorithm in the form an implementation needs. P_k denotes continuous piecewise polynomials of degree k on the mesh, P_k^{disc} the discontinuous ones, and a bubble a polynomial on one element that

vanishes on a prescribed part of its boundary; U_h/Q_h denotes the pair with displacement space U_h and pressure space Q_h .

4.1 The pair

Let $U_h \subset \mathcal{U}$ carry the motion and $Q_h \subset \mathcal{Q}$ carry both $\bar{\theta}$ and $\bar{p}$. Two requirements fix the choice.

Same space for both scalars. $\bar{\theta}$ and $\bar{p}$ are placed in one space Q_h . This is what makes the constraint term vanish identically and the functional reduce to (19); it is the only thing that reduction needs. Given it, and a quadratic W^{vol} (§3.5), (33)₂ reads $\bar{p} = \kappa \bar{\theta}$ pointwise rather than in a projected sense, and the composition $\bar{p} = \kappa \mathcal{P}_h(\theta)$ is well defined. Splitting the two scalars across different spaces makes the constitutive relation hold only after a projection, determines one field only through its projection onto the other's space, and can leave the discrete system singular.

Inf-sup stability, in three dimensions. By (34) the pressure–displacement pairing is $\int \varpi \mathbf{F}^{-\top} : \text{Grad } \boldsymbol{\psi}$ with $\varpi = \theta'(J) J \bar{p}$, and not $\int \bar{p} \mathbf{F}^{-\top} : \text{Grad } \boldsymbol{\psi}$; the two coincide at $\mathbf{F} = \mathbf{I}$ and nowhere else. The normalization (13) sets $\theta'(1) = 1$, so $\varpi = \bar{p}$ there and the pairing reduces to the divergence pairing $\int \bar{p} \text{Div } \boldsymbol{\psi} \, dV$ — in both instances of Table 2, which is what makes the stability results of §9.5 a property of the pair rather than of the material. That pairing is the one of incompressible linear elasticity, in which the pressure is the multiplier of the constraint $\text{Div } \boldsymbol{\varphi} = 0$, and its discrete inf-sup condition,

$$\inf_{q \in Q_h} \sup_{\mathbf{v} \in U_h} \frac{\int_{\Omega} q \, \text{Div } \mathbf{v} \, dV}{|\mathbf{v}|_{H^1} \|q\|_{L^2}} \geq \beta_h, \quad (37)$$

is the condition on the pair. The admissible pairs are therefore those for which β_h in (37) is bounded below uniformly in h , and stability must also hold on the target range of κ/μ [6].

The pair has been used before in large-deformation hyperelasticity: Chamberland et al. [8], comparing discretizations of a nearly incompressible Mooney–Rivlin material against manufactured solutions on tetrahedral and hexahedral meshes, recommend it and its hexahedral counterpart for nearly incompressible large-deformation problems for accuracy, run time and memory, and observe on some manufactured solutions a larger pressure error than that of Taylor–Hood, which they attribute to a smaller inf-sup constant of the richer pressure space (the constants measured in §9.5 are 0.2968 for this pair and 0.2216 for Taylor–Hood on the Freudenthal meshes). Their formulation is the two-field displacement–pressure form, whose pressure equation $\int_{\Omega_0} (J - 1) q \, dV - \int_{\Omega_0} \kappa^{-1} p q \, dV = 0$ for all q in the discontinuous space Q_h gives $p = \kappa \mathcal{P}_h(J - 1)$ element by element; in elasticity their element is therefore the element of this note with $\theta = J - 1$ and a volumetric energy quadratic in θ . They also report that on meshes of curved elements a pressure linear in the coordinates of the physical element is more accurate than one linear in the coordinates of the parent element. The element of this note takes the pressure linear in the parent coordinates $\boldsymbol{\xi}$ (§4.4); on an element with straight edges these are affine functions of the reference coordinates $\mathbf{X}$ and the two choices coincide, and they differ only on elements with curved edges, such as those on the lateral surface of the Taylor bar (§11). The effect of that choice on the element of this note is not measured. No use of the pair in finite-deformation plasticity is known to the author.

Stability of a pair depends on the spatial dimension. The conforming Crouzeix–Raviart pair [11] is, in two dimensions, P_2 enriched with a cubic interior bubble over P_1^{disc} ; its three-dimensional counterpart below carries a quartic interior bubble and cubic face bubbles, and the quadrature degree the element needs follows from the quartic. Unenriched P_2/P_1^{disc} is stable in two dimensions only on special meshes. In three dimensions Boffi et al. [6, §8.7] state that a constant discontinuous pressure is not controlled by P_2 : “to control piecewise constant pressure, we need some degrees of freedom on the faces,” and the quadratic space has none. The stable three-dimensional choices with a discontinuous pressure therefore require bubble enrichment of the displacement space, and the enrichment has two parts with two functions: face bubbles control the constant pressure, and an interior bubble lifts it to a linear one [6, Example 8.7.2]. Table 3 lists the candidates, and §9.5 measures every one of them.

Displacement	Pressure	3D stable	Pressure elimination
P_1	P_0	no	local
P_2	P_0	no	local
$P_2 \oplus b_{\text{face}}$	P_0	yes	local
P_2	P_1^{disc}	no	local
$P_2 \oplus b_{\text{int}}$	P_1^{disc}	no	local
$P_2 \oplus b_{\text{int}} \oplus b_{\text{face}}$	P_1^{disc}	yes	local
P_2	P_1 continuous	yes	global

Table 3: Displacement–pressure pairs on tetrahedra, with three-dimensional stability as Boffi et al. [6, §8.7, Example 8.7.2] give it. Local elimination requires a discontinuous pressure; three-dimensional stability with a discontinuous pressure requires face bubbles, and a linear pressure requires an interior bubble as well.

The pair this formulation requires is P_2 enriched with face and interior bubbles over P_1^{disc} . The two enrichments are different kinds of object. The interior bubble vanishes on the whole element boundary: it is never shared, never reached by a boundary condition, and condenses away, adding no global unknown. A face bubble vanishes on three faces and not on its own, so the two elements meeting there share it; it is a global unknown and does not condense. The enrichment therefore adds global degrees of freedom, which §9.3 counts. The pressure elimination of §4.2 nevertheless stays element by element, because it rests on Q_h being discontinuous and not on the displacement space.

The linear pressure is not forced by the stability results. $P_2 \oplus b_{\text{face}}/P_0$ is also stable (§9.5), at a higher constant and with the same number of global unknowns, since the interior bubble condenses. The count of §9.2 applies to the unenriched displacement space, and §9.4 shows that its modes do not persist after assembly. What separates the two pressure spaces is the claim of §6: a constant pressure leaves in $\ker \mathbf{G}_h$ the element-scale dilatational modes that a plastic tangent can free, and a linear one does not. §9.7 measures both pairs, and that measurement is the reason for the linear pressure.

Both that pair and the Taylor–Hood pair — continuous P_2 over continuous P_1 , whose stability in three dimensions is established [6] and which serves as the positive control of the measurements — are measured in §9.5, and every other element-local pair decays. The difference between the two stable pairs is not stability but where the coupling is — a global pressure that cannot be eliminated element by element against global face bubbles that cannot

be condensed — and §13 quantifies it.

4.2 Local elimination

Because Q_h is discontinuous, the mass matrix $[M]_{mn} = \int_{\Omega} \chi_m \chi_n \, dV$, with $\{\chi_m\}$ a basis of Q_h , is block diagonal by element and constant in the reference configuration. One factorization at setup serves every Newton iteration of every step, and

$$M\{\bar{\theta}\} = \int_{\Omega} \chi_m \theta(\varphi) \, dV, \quad \bar{p} = \kappa \bar{\theta}, \quad (38)$$

is an element-local solve, with $\theta = \log J$ for a Hencky material and $\theta = J - 1$ for the Simo–Hughes one (Table 2). The second relation in (38) holds pointwise, rather than only after projection, precisely because W^{vol} is quadratic; for a material where it is not, it is replaced by a second projection with the same factorization (§3.6). Substituting $\bar{p}(\varphi)$ into (33)₁ leaves a displacement-sized nonlinear problem whose tangent retains the sparsity of the displacement mesh — the auxiliary fields never couple two elements. This is the property that keeps the composite tetrahedron’s tangent as sparse as a displacement element’s, and it is preserved here.

Remark 4. The locality is a consequence of the discontinuous pressure, not of the Hu–Washizu structure. A conforming auxiliary field — one requiring inter-element continuity — would make $\bar{p}(\varphi)$ a global function of φ , and every tangent–vector product would then require a global solve with its mass matrix, which the composite tetrahedron does not require. That is the main reason the pressure is kept discontinuous.

4.3 The element

The displacement space of the pair is carried by a fifteen-node tetrahedron, **TETRA15** in the naming of the Exodus mesh format, with three displacement unknowns at each node. This section defines it in full.

Nodes. With $\lambda_1, \dots, \lambda_4$ the barycentric coordinates of the tetrahedron, the fifteen nodes are the four vertices (nodes 1–4); the six edge midpoints (nodes 5–10), on the edges 1–2, 2–3, 3–1, 1–4, 2–4, 3–4 in that order; the four face centroids (nodes 11–14), on the faces with vertices 1–3–2, 2–3–4, 1–4–3 and 1–2–4 in that order, which is the side-node convention of the Exodus library (node 14 on side 1, node 12 on side 2, node 13 on side 3, node 11 on side 4); and the centroid (node 15). The first ten are the nodes of the quadratic tetrahedron **TET10** in the same order, so a **TETRA10** mesh is converted by adding one node at the centroid of every face and of every element; for a mesh with straight edges the added nodes are the averages of the face and element vertices. Each face is a seven-node triangle: three vertices, three edge midpoints and the centroid.

The space. The displacement space on one element is

$$[P_2 \oplus \text{span}\{b_{F_1}, \dots, b_{F_4}\} \oplus \text{span}\{b_T\}]^3, \quad b_F = 27 \lambda_a \lambda_b \lambda_c, \quad b_T = 256 \lambda_1 \lambda_2 \lambda_3 \lambda_4, \quad (39)$$

with P_2 the ten quadratic Lagrange functions of the tetrahedron, b_F the cubic bubble of face F with vertices a, b, c , which vanishes on the three other faces and equals one at the centroid of

F , and b_T the quartic interior bubble, which vanishes on the whole boundary and equals one at the centroid. The space has dimension fifteen per component and contains the complete quadratic polynomials, so the element is conforming and its interpolation error is that of the quadratic tetrahedron.

The nodal basis. The functions of (39) in the order listed — the ten Lagrange functions of TET10, then $b_{F_1}, \dots, b_{F_4}$ in the face order of the nodes, then b_T — form the hierarchical basis $\{\hat{N}_j\}$, so called because the functions of the lower-degree element remain in it and the bubbles are appended. It is not nodal: the Lagrange functions are nonzero at the face centroids and at the centroid. The nodal basis $\{N_a\}$, with $N_a(\mathbf{x}_b) = \delta_{ab}$ at the fifteen nodes, is obtained from it by the constant matrix

$$N_a = \sum_j A_{aj} \hat{N}_j, \quad \mathbf{A} = (\mathbf{V}^{-1})^\top, \quad V_{bj} = \hat{N}_j(\mathbf{x}_b), \quad (40)$$

with $\mathbf{V}$ the values of the hierarchical functions at the nodes, a 15×15 matrix with rational entries that is inverted exactly once. The gradients transform by the same matrix. Every unknown is then the displacement of a node, so a code that numbers unknowns by node treats the element like any Lagrange element: the four face-centroid unknowns are shared between the two elements of a face and the centroid unknown belongs to one element. §9.8 measures the two bases; the nodal one is used.

Quadrature. The element is integrated with the fourteen-point rule of Keast [17], exact for polynomials of degree five, with positive weights. Its points are three orbits of the symmetry group of the tetrahedron: two orbits of four points with barycentric coordinates $(1 - 3a, a, a, a)$ and its permutations, and one orbit of six points with coordinates $(b, b, \frac{1}{2} - b, \frac{1}{2} - b)$ and its permutations,

$$\begin{aligned} a_1 &= 0.0927352503108912, & w_1 &= 0.01224884051939366, \\ a_2 &= 0.3108859192633006, & w_2 &= 0.01878132095300264, \\ b &= 0.4544962958743503, & w_3 &= 0.007091003462846911, \end{aligned} \quad (41)$$

with the weights normalized to the reference volume $1/6$. The gradient of b_T is cubic and the stiffness integrand of its block is of degree six, which the rule integrates with an error of 35% on that block (§9.8); the rules of degree two and three (four and five points) leave the element twenty-one zero-energy modes beyond the six rigid-body modes and cannot be used. The same rule integrates the mass matrix, the projection (38) and the residual.

Mass. The consistent mass matrix in the nodal basis has condition number 94 on the reference element. Its row sums, $m_a = \rho \int_e N_a dV$, are positive at every node because the nodal basis is a partition of unity with positive integrals: as fractions of the element mass they are 0.020 at each vertex, 0.038 at each edge midpoint, 0.096 at each face centroid and 0.305 at the centroid. The row-sum lumped mass is therefore admissible for explicit time integration (§7), which it is not for TET10, whose vertex row sums are negative.

4.4 Element algorithm

This section states what an implementation computes on one element, with the nodal displacements $\mathbf{u}_a$ of its fifteen nodes given. The pressure space Q_h on the element is spanned by $\chi_m(\boldsymbol{\xi})$, $m = 1, \dots, 4$, with $\chi = (1, \xi_1, \xi_2, \xi_3)$ in the reference coordinates $\boldsymbol{\xi}$ of the element; the constant projection, the element TET15-P0, which §10 also measures, has $\chi = (1)$. The mass matrix of the projection is taken with the reference-configuration volume element, so it does not change with the deformation. Two passes over the quadrature points $\boldsymbol{\xi}_q$, with weights w_q and reference Jacobian determinants j_q , compute the residual.

First pass: the projected volumetric strain. At each quadrature point, $\mathbf{F}_q = \mathbf{I} + \sum_a \mathbf{u}_a \otimes \text{Grad } N_a(\boldsymbol{\xi}_q)$, $J_q = \det \mathbf{F}_q$ and $\theta_q = \theta(J_q)$. Then

$$M_{mn} = \sum_q w_q j_q \chi_m(\boldsymbol{\xi}_q) \chi_n(\boldsymbol{\xi}_q), \quad b_m = \sum_q w_q j_q \chi_m(\boldsymbol{\xi}_q) \theta_q, \quad \bar{\theta} = \mathbf{M}^{-1} \mathbf{b}, \quad (42)$$

a 4×4 solve, and the projected strain and pressure at each point are $\bar{\theta}_q = \sum_m \chi_m(\boldsymbol{\xi}_q) \bar{\theta}_m$ and $\bar{p}_q = \kappa \bar{\theta}_q$.

Second pass: the stress and the residual. At each quadrature point the material supplies the isochoric first Piola–Kirchhoff stress $\mathbf{P}_q^{\text{iso}}$ from $\mathbf{F}_q$ and its internal variables, with its own return map, and the element adds the volumetric part with the projected pressure:

$$\mathbf{P}_q = \mathbf{P}_q^{\text{iso}} + \bar{p}_q \mathbf{g}_q, \quad \mathbf{g}_q := \left. \frac{\partial \theta}{\partial \mathbf{F}} \right|_q = \theta'(J_q) J_q \mathbf{F}_q^{-\top}, \quad \mathbf{R}_a = \sum_q w_q j_q \mathbf{P}_q \text{Grad } N_a(\boldsymbol{\xi}_q). \quad (43)$$

The residual is the pointwise one with the projected pressure in place of $\kappa \theta_q$; it is the derivative of the element energy $\sum_q w_q j_q [W^{\text{iso}}(\bar{\mathbf{F}}_q) + \frac{\kappa}{2} \bar{\theta}_q^2]$ with respect to the nodal displacements, which §4.5 verifies by differences.

The tangent. The derivative of (43) along a nodal displacement increment $\mathbf{v}$ with gradient $\text{Grad } \mathbf{v}_q$ at the quadrature points has three parts: the isochoric material tangent of the model, the geometric term of the volumetric stress, and the coupling through the projection,

$$\begin{aligned} \delta \mathbf{P}_q &= \left. \frac{\partial \mathbf{P}^{\text{iso}}}{\partial \mathbf{F}} \right|_q : \text{Grad } \mathbf{v}_q + \bar{p}_q \left. \frac{\partial \mathbf{g}}{\partial \mathbf{F}} \right|_q : \text{Grad } \mathbf{v}_q + \kappa \delta \bar{\theta}_q \mathbf{g}_q, \\ \delta \bar{\theta} &= \mathbf{M}^{-1} \sum_q w_q j_q \chi(\boldsymbol{\xi}_q) (\mathbf{g}_q : \text{Grad } \mathbf{v}_q), \quad \delta \bar{\theta}_q = \sum_m \chi_m(\boldsymbol{\xi}_q) \delta \bar{\theta}_m, \end{aligned} \quad (44)$$

with, for $\phi(J) := \theta'(J)J$ and ϕ' its derivative,

$$\left. \frac{\partial \mathbf{g}}{\partial \mathbf{F}} \right|_q : \delta \mathbf{F} = \phi'(J) J (\mathbf{F}^{-\top} : \delta \mathbf{F}) \mathbf{F}^{-\top} - \phi(J) \mathbf{F}^{-\top} \delta \mathbf{F}^\top \mathbf{F}^{-\top}. \quad (45)$$

The third term couples the quadrature points of the element: in matrix form it adds $\kappa \mathbf{D} \mathbf{M}^{-1} \mathbf{D}^\top$ to the sum of the quadrature-point matrices, with $\mathbf{D} = \sum_q w_q j_q (\text{Grad } N \cdot \mathbf{g}_q) \chi(\boldsymbol{\xi}_q)^\top$ the matrix that maps nodal displacements to the right-hand side of (42). The tangent is symmetric, since it is the second derivative of the element energy. The matrix-free tangent–vector product is (44) applied to $\mathbf{v}$ with the two passes; the assembled tangent, its diagonal and its nodal blocks are formed from the same three terms.

What the material supplies. A material admissible for the element provides, at a quadrature point with its internal variables: the volumetric strain measure $\theta(J)$ with $\theta'(J)$ and $\theta''(J)$; the bulk modulus κ ; the isochoric energy $W^{\text{iso}}(\bar{\mathbf{F}})$; the isochoric first Piola–Kirchhoff stress, together with the update of the internal variables it entails; and the isochoric tangent $\partial \mathbf{P}^{\text{iso}} / \partial \mathbf{F}$. For the Simo–Hughes J_2 model, $\theta = J - 1$, $\theta' = 1$, $\theta'' = 0$, $\kappa = \lambda + 2\mu/3$ from the Lamé constants, and the isochoric stress and tangent are those of BOX 9.1 of Simo and Hughes [25] with the volumetric term omitted; the return map on $\bar{\mathbf{b}}^e$ and the internal variables $\mathbf{F}^p$ and $\bar{\mathbf{e}}^p$ are unchanged. For a material without an accessible split, the passes and the stress are those of §3.6, with the material called at the modified deformation gradient (25).

Assembly. Because the first pass needs every quadrature point of the element before the second can begin, the element is assembled as a unit, not quadrature point by quadrature point. Internal variables are stored per quadrature point as for any element. Traction on a face are integrated over the seven-node triangle with the triangle rule of the same degree as the element rule. The element runs under the quasi-static, the implicit and the explicit integrators of the code without change.

4.5 Consistency checks

An implementation of §4.4 is verified by the following checks, which the implementation in Carina passes (`test/projected-element.jl`, forty-six checks) on a cube of TET15 elements with the J_2 model in its elastic range at a converged nonlinear state, so that the geometric and coupling terms of the tangent are nonzero, for both the linear and the constant projection: the residual equals the derivative of the assembled energy by central differences, and the assembled tangent the derivative of the residual, to the truncation error of the differences; the matrix-free action, the diagonal and the nodal-block kernels reproduce the assembled tangent to 10^{-12} ; an affine displacement field is reproduced exactly (patch test, 10^{-10}), because a constant J is its own projection and the projected element must then agree with the pointwise one; and Newton converges quadratically (relative residual 1, 8.6×10^{-3} , 8×10^{-7} , 2×10^{-14}).

The general form of §3.6 passes the same checks (`test/projected-general.jl`, 129 checks) with the neo-Hookean material, which has no accessible split in the material library, for $\theta = \log J$ and $\theta = J - 1$ and for both projections, on the CPU and on a GPU, where the Newton history and the displacements agree with the CPU’s within 10^{-8} . Its reduction to the split form is checked on the same cube with the J_2 model forced through the general form with $\theta = J - 1$: the Newton history, the displacements, the internal variables, the energy, the residual and the assembled tangent agree with the split form’s to between 10^{-15} and 10^{-13} relative, in the elastic range and in a plastic step with an equivalent plastic strain of 0.021. The formulas themselves were checked on one distorted element before the kernels were written (`prototype/general.jl`): the hand-derived residual against the gradient of the energy by automatic differentiation to 4×10^{-16} , the closed-form tangent against its Hessian to 1.2×10^{-14} , and the three reductions of §3.6 to rounding. The same prototype found that the tangent of the J_2 model in the material library, the consistent tangent of BOX 9.2 of Simo and Hughes [25], divided its coefficient β_2 by μ instead of the effective modulus $\bar{\mu} = \mu \text{tr}(\bar{\mathbf{b}}_{\text{tr}}^e)/3$; the two agree at small strain and the term vanishes without hardening, which is why the library’s own checks had not seen it, and the symmetric part of the tangent was off by 5×10^{-2} at an equivalent plastic strain of 0.6 reached in one increment. With the correction the symmetric part agrees with the central-difference

Jacobian of the stress to 10^{-10} ; what remains, 4×10^{-3} to 5×10^{-3} , is the antisymmetric part that BOX 9.2 drops by construction, so Newton converges linearly, not quadratically, on steps that yield with this model.

5 Plasticity

This section develops the logarithmic instance of Table 2; the Simo–Hughes column differs here and nowhere else (§3.7), and Remark 5 states what the implemented instance, Simo–Hughes, entails.

The deviatoric response is evaluated at quadrature points of the displacement element, on $\mathbf{e}(\boldsymbol{\varphi})$ computed there. The plastic strain is carried as a logarithmic strain $\mathbf{e}^p$ and subtracted additively — the construction of Miehe et al. [18], in which a material geometric pre-processing $\mathbf{C} \mapsto \boldsymbol{\epsilon}$ and a post-processing through $\mathbb{L}(\mathbf{C})$ precede and follow an unmodified small-strain algorithm. This is a different constitutive model from the multiplicative $\mathbf{F} = \mathbf{F}^e \mathbf{F}^p$ plasticity of Simo [24], which also preserves the small-strain return map but applies it to the elastic logarithmic strain of a trial elastic state with an exponential update; the two coincide under proportional loading with fixed principal axes and differ when the axes rotate. Because the constitutive law is linear in logarithmic strain, the J_2 return map is the small-strain radial return — the return map that scales the trial deviatoric stress toward the yield surface along its own direction — unchanged:

$$\begin{aligned} \mathbf{e}^{\text{tr}} &:= \mathbf{e}(\boldsymbol{\varphi}) - \mathbf{e}_n^p, & \mathbf{s}^{\text{tr}} &:= 2\mu \mathbf{e}^{\text{tr}}, \\ \Delta\gamma : \quad \|\mathbf{s}^{\text{tr}}\| - 2\mu \Delta\gamma - \sqrt{2/3} \left[\sigma_y + \mathcal{H}'(\bar{\epsilon}_n^p + \sqrt{2/3} \Delta\gamma) \right] &= 0 & \text{if } f^{\text{tr}} > 0, \quad \text{else } \Delta\gamma := 0, \\ \mathbf{e}_{n+1}^p &:= \mathbf{e}_n^p + \Delta\gamma \mathbf{n}, & \mathbf{n} &:= \mathbf{s}^{\text{tr}} / \|\mathbf{s}^{\text{tr}}\|, & \bar{\epsilon}_{n+1}^p &:= \bar{\epsilon}_n^p + \sqrt{2/3} \Delta\gamma, \\ \mathbf{s} &:= 2\mu (\mathbf{e}(\boldsymbol{\varphi}) - \mathbf{e}_{n+1}^p). \end{aligned} \tag{46}$$

Here σ_y is the initial yield stress, $\bar{\epsilon}^p$ the equivalent plastic strain, $\mathcal{H}'$ the hardening function (the increase of the yield stress with $\bar{\epsilon}^p$) and $\mathcal{H}''$ its derivative, the hardening modulus, and $f^{\text{tr}} := \|\mathbf{s}^{\text{tr}}\| - \sqrt{2/3} [\sigma_y + \mathcal{H}'(\bar{\epsilon}_n^p)]$ the trial yield function. The consistency condition is a scalar equation in $\Delta\gamma$. For linear hardening it closes in one step, $\Delta\gamma = f^{\text{tr}} / (2\mu + \frac{2}{3}\mathcal{H}'')$, which is the form both codes implement; for nonlinear hardening it is a scalar Newton iteration [25, BOX 3.2]. Internal variables are stored at quadrature points and are never projected, so no interpolation error enters the history. Frame-indifferent interpolation, when tensorial internal variables must be transferred, follows Mota et al. [19].

Three properties follow. They are the motivation for the formulation rather than incidental consequences of it.

The constitutive law is untouched. All finite-deformation content is absorbed into two material-independent geometric transforms: the forward map $\mathbf{C} \mapsto \boldsymbol{\epsilon}$ and the tangent map $\mathbb{L}(\mathbf{C})$ [18]. Between them is an unmodified small-strain algorithm. Any model already written for small strain — J_2 , kinematic hardening, viscoplasticity, damage — is usable without reformulation. This is the practical content of preferring a Hu–Washizu structure over one built around a particular stress–strain pairing. This property belongs to the logarithmic measure

and not to the Hu–Washizu structure, so it is the one place where §3.5 does not apply: an admissible material in the sense of (20) need not have it, and the Simo–Hughes model both codes implement does not (Remark 5).

Strain and stress never cross discretizations. $\mathbf{e}$ is evaluated and $\mathbf{s}$ returned at the same quadrature point of the same element. There is no projection between the two, so the interpolation error between two discrete spaces that corrupts the equivalent plastic strain in stabilized formulations has no mechanism here.

No stabilization enters the deviator. There is no subgrid term, no penalty and no tuned coefficient anywhere in the formulation, so nothing of the form (1) can reach the return map.

Remark 5 (The instance implemented). The instance implemented in Carina, and the one every nonlinear result of §10 uses, is the Simo–Hughes J_2 model [25, BOX 9.1]: the multiplicative decomposition $\mathbf{F} = \mathbf{F}^e \mathbf{F}^p$ with the volumetric–isochoric split $W = \frac{\kappa}{2}(J-1)^2 + \frac{\mu}{2}(\text{tr } \bar{\mathbf{b}}^e - 3)$, and the radial return taken on the isochoric elastic left Cauchy–Green tensor $\bar{\mathbf{b}}^e$ rather than on the logarithmic strain. It satisfies (20) exactly — $\det \bar{\mathbf{b}}^e = 1$ makes the isochoric term independent of J by construction — and its W^{vol} is quadratic in $J - 1$, so the element of §4 uses it in the variable $\theta = J - 1$ with no change other than the omission of the volumetric term from the stress the model returns. The two instances share the mixed treatment of the volume and differ in the plasticity model: additive plasticity in logarithmic strain in (46), multiplicative plasticity with the return map on $\bar{\mathbf{b}}^e$ in the implemented instance; the two coincide under proportional loading with fixed principal axes and differ when the axes rotate. The property that a small-strain model is usable without reformulation belongs to the logarithmic instance alone; the Simo–Hughes return map is specific to its own kinematics. The smoothness of the equivalent plastic strain is a question of whether strain and stress cross discretizations, which §4.4 settles for either instance, and it is to be measured with the implemented one (§12).

6 The claim

This section states the claim of the note: which hypothesis of Brezzi’s theorem a plastic tangent can violate, and what distinguishes pairs under it. The second hypothesis is coercivity of the deviatoric form on $\ker \mathbf{G}_h$ (§1).

In elasticity the second hypothesis holds for every conforming pair: by Korn’s inequality, which bounds $\int |\text{Grad } \mathbf{v}|^2 dV$ by a multiple of $\int |\text{dev } \boldsymbol{\epsilon}(\mathbf{v})|^2 dV$ for $\mathbf{v} \in H_0^1$, the deviatoric form is coercive on all of U_h , so it is coercive on $\ker \mathbf{G}_h$, the set of discrete displacements whose divergence has zero projection onto Q_h (§9.1), and no elastic soft mode is a failure of the theorem. A count of exact zero-energy modes on one element is therefore not an argument against a constant pressure in elasticity; §9.4 shows that those modes do not persist after assembly.

Under plastic loading the consistent J_2 tangent, the derivative of the stress returned by (46) with respect to the strain, is

$$\mathbb{C}^{\text{ep}} = 2\mu \left(\mathbb{I}^{\text{dev}} - \eta \mathbf{n} \otimes \mathbf{n} \right), \quad \mathbb{I}^{\text{dev}} := \mathbb{I} - \frac{1}{3} \mathbf{I} \otimes \mathbf{I}, \quad \mathbf{n} := \mathbf{s}^{\text{tr}} / \|\mathbf{s}^{\text{tr}}\|, \quad \eta = \frac{2\mu}{2\mu + \frac{2}{3}\mathcal{H}''}, \quad (47)$$

with $\eta = 1$ for perfect plasticity: the shear stiffness along the flow direction is removed.

Read as a constant, the second hypothesis cannot separate pairs. For every $\mathbf{v} \in H_0^1(\Omega)$, integration by parts gives $\int |\text{dev } \boldsymbol{\epsilon}|^2 dV = \frac{1}{2} \int |\text{Grad } \mathbf{v}|^2 dV + \frac{1}{6} \int (\text{Div } \mathbf{v})^2 dV$, and $(\mathbf{n} : \boldsymbol{\epsilon})^2 \leq |\text{dev } \boldsymbol{\epsilon}|^2$ since $\mathbf{n}$ is a unit deviator, so

$$\frac{1}{2\mu} \int \boldsymbol{\epsilon} : \mathbb{C}^{\text{ep}} : \boldsymbol{\epsilon} dV \geq (1 - \eta) \int |\text{dev } \boldsymbol{\epsilon}|^2 dV \geq \frac{1 - \eta}{2} \int |\text{Grad } \mathbf{v}|^2 dV \quad (48)$$

for every conforming pair and every member of its kernel, isochoric or not. For $\eta < 1$ the coercivity constant on $\ker \mathbf{G}_h$ is therefore bounded below uniformly in h and in the pair. For $\eta = 1$ the continuum operator loses strong ellipticity, positivity of the energy of every displacement jump across a plane: a shear band, a displacement jump $\mathbf{e}_1$ across a plane with normal $\mathbf{e}_2$, has $\boldsymbol{\epsilon} \propto \text{sym}(\mathbf{e}_1 \otimes \mathbf{e}_2)$, to which (47) with that $\mathbf{n}$ assigns zero energy, so the constant tends to zero under refinement for every pair as the discrete space resolves narrower bands (§9.7). Neither limit says anything about the pair.

What distinguishes the pairs is what the soft directions are. A member of $\ker \mathbf{G}_h$ that is nearly isochoric is a shear band the continuum problem also has, and its low energy is physical. A member with pointwise nonzero dilatation of zero element mean has the volumetric energy $\kappa \int (\text{Div } \mathbf{v})^2 dV$ in the continuum and none in the element, so once the tangent removes its shear energy the element assigns it zero energy. That is the soft mode reported for constant-pressure elements, and whether $\ker \mathbf{G}_h$ contains such members is a property of the pair. For a pair satisfying the inf-sup condition, $\ker \mathbf{G}_h$ approximates the space of isochoric displacements at the rate of U_h [6, Prop. 5.1.3], so the dilatation carried by any of its members vanishes with h . For an element-constant pressure on a quadratic displacement, $\ker \mathbf{G}_h$ contains displacements whose dilatation is nonzero pointwise and averages to zero over each element, and their dilatation does not vanish with h .

The claim is therefore the following. Under a plastic tangent, every soft mode of the stable pair with a linear discontinuous pressure carries a dilatation $\int (\text{Div } \mathbf{v})^2 dV / \int |\text{Grad } \mathbf{v}|^2 dV$ that vanishes under refinement, and a constant pressure has soft modes whose dilatation does not; the soft modes reported for constant-pressure elements are the latter, and no stabilization is needed where they are absent. Both operators are linear at a fixed plastic state, so the claim is tested without a nonlinear solve, and §9.7 tests it. If the dilatation of the soft modes does not separate the pairs, the formulation has no advantage over the constant-pressure elements it is meant to replace.

7 Dynamics

The spatial formulation is unchanged; the mass matrix enters only the displacement block. Two implementation choices are specified.

Lumping. Explicit integration divides by a diagonal mass, so every lumped mass must be positive. Row-sum lumping assigns to each node the row sum of the consistent mass matrix, $m_a = \rho \int N_a dV$, which for a basis whose functions sum to one adds up to the element mass. On the quadratic tetrahedron TET10 the vertex functions have negative integrals, so the row-sum mass is negative at the vertices and cannot be used. In the nodal basis of the enriched element (§4.3) every row sum is positive: 0.020 of the element mass at each vertex, 0.038 at each edge midpoint, 0.096 at each face centroid and 0.305 at the centroid, on the reference element and on

the distorted and Kuhn elements of §9.8. The row-sum lumped mass is used; it conserves the momentum of a rigid translation exactly, because the functions sum to one. The alternative is diagonal-scaling lumping [15] (HRZ: Hinton, Rock and Zienkiewicz), which takes the diagonal of the consistent mass, positive for any basis, and scales it to preserve the element mass. In the nodal basis HRZ also conserves momentum. In the hierarchical basis it does not: a rigid translation has zero coefficient on every bubble unknown, so the mass HRZ assigns to the bubbles, 51% of the element mass on the reference tetrahedron, carries no momentum, unless the scaling runs over the ten Lagrange functions alone with each bubble keeping its consistent diagonal (§9.8).

Critical time step. From the element eigenvalue bound at $\kappa/\mu = 2$ (§9.8), the Courant number $c_p \Delta t/h$, with $c_p = \sqrt{(\kappa + 4\mu/3)/\rho}$ the dilatational wave speed, ρ the density, Δt the stable step and h the element edge length, is given in Table 12 for the four-node linear tetrahedron (TET4), for TET10 under HRZ, and for the enriched element in both bases. The quadratic element reduces the stable step by a factor of 2.4 against the linear one on the same mesh, and the enrichment a further 1.4 with the nodal basis; the ordering is the same on the distorted element and on a Kuhn tetrahedron of the Freudenthal mesh (§9.1). A P_2 mesh resolves a given feature with fewer elements; on the Taylor bar (§11) the smaller element count more than compensates for the smaller step. These Courant numbers hold on the undeformed element: on a deformed mesh the smallest distance between two nodes does not bound the critical step, which is then computed from the largest eigenvalue of $M_L^{-1}K$ (§11).

8 Optional five-field variant

This section gives the extension to an independent deviatoric strain field, for non-local regularization or for internal variables requiring their own discrete representation: it introduces $\bar{\epsilon}$ and $\bar{s}$ in the same discontinuous space, of the same degree as Q_h . Sharing the space keeps the return map at shared quadrature points, keeps every elimination element-local, leaves the discrete system non-singular, and — for a linear deviatoric law, and only then — makes the constitutive relation hold pointwise rather than after a projection; with plasticity $\bar{s} = \mathcal{P}_h s(\bar{\epsilon})$ is projected. Placing $\bar{\epsilon}$ and $\bar{s}$ in different spaces loses each of these properties. Remark 2 also applies: with an independent $\bar{s}$, the residual must be assembled from (35) and not from (36).

9 Numerical results

Every measurement below is on a linear operator, and all but those of §9.6 are at $\mathbf{F} = \mathbf{I}$: the condensed operator (51) for the count and the locking sweep, the H^1 seminorm for the stability sweep, and the deviatoric form under a plastic tangent for the soft-mode measurements. None needs a solver, a material model, or a mesh reader. The scripts are in `prototype/`, each reports its own result, and `checks.jl` verifies the quadrature, the bubbles, and the conformity of the enriched spaces against values known in advance.

9.1 Method

Meshes are generated rather than read, so the refinement sequence is exact. The unit cube is divided into N^3 cubes and each cube into six tetrahedra by the Freudenthal (Kuhn) subdivision

— the six permutations (a, b, c) giving the Kuhn tetrahedra $\{\mathbf{0}, \mathbf{e}_a, \mathbf{e}_a + \mathbf{e}_b, (1, 1, 1)\}$. That decomposition is conforming across shared cube faces, which the five-tetrahedron split is not. Quadratic meshes add edge midpoints in the node order the reference element expects.

For the stability test the displacement is fixed on the entire boundary and three operators are assembled,

$$K_{ij} = \int \text{Grad } N_i : \text{Grad } N_j \, dV, \quad G_{mi} = \int \chi_m \text{Div } N_i \, dV, \quad M_{mn} = \int \chi_m \chi_n \, dV, \quad (49)$$

with N_i the displacement basis functions, χ_m the pressure basis functions, and K the H^1 seminorm, the norm on the displacement in (37); $\ker \mathbf{G}$ is the discretely isochoric subspace, the set of displacements whose divergence has zero projection onto Q_h . The discrete inf-sup constant is $\beta_h = \sqrt{\lambda_{\min}}$ over the nonzero spectrum of

$$\mathbf{G} \mathbf{K}^{-1} \mathbf{G}^\top q = \lambda \mathbf{M} q. \quad (50)$$

Boundary data fixed everywhere always leaves one zero mode, the hydrostatic one, because $\int \text{Div } \mathbf{v} = 0$ for $\mathbf{v} \in H_0^1$; any zero beyond that is a spurious pressure mode.

9.2 Spurious modes on one element

For linear response the count of zero-energy modes is a rank count on a single element. Condensing the three fields gives

$$\mathbf{K}_e = \mathbf{K}^{\text{dev}} + \mathbf{K}^{\text{vol}}, \quad \mathbf{K}^{\text{vol}} := \kappa \mathbf{G}^\top \mathbf{M}^{-1} \mathbf{G}, \quad M_{mn} = \int_e \chi_m \chi_n \, dV, \quad G_{mi} = \int_e \chi_m (\text{tr } \mathbf{B})_i \, dV, \quad (51)$$

with $\mathbf{B}$ the strain–displacement operator of the element and $m = \dim Q_h$ on one element. $\mathbf{K}^{\text{dev}}$ annihilates every mode with zero deviatoric strain at quadrature: the six rigid-body modes, plus the volumetric directions the displacement space can produce. The second term has rank equal to the number of those directions the pressure space resists, at most m . Anything left over carries neither deviatoric nor volumetric energy and is exactly zero-energy:

$$n_{\text{spurious}} = \dim \ker \mathbf{K}^{\text{dev}} - 6 - \text{rank } \mathbf{K}^{\text{vol}}. \quad (52)$$

For the quadratic tetrahedron $\dim \ker \mathbf{K}^{\text{dev}} = 10$: six rigid modes and four volumetric directions, the latter because $\text{tr } \boldsymbol{\epsilon}(\boldsymbol{\varphi})$ is a P_1 field. Table 4 reports (52) against the measured spectrum.

Pressure space	m	$\dim \ker \mathbf{K}^{\text{dev}}$	$\text{rank } \mathbf{K}^{\text{vol}}$	predicted	observed
P_0 constant	1	10	1	3	3
P_1 discontinuous	4	10	4	0	0
P_2 discontinuous	10	10	4	0	0

Table 4: Spurious zero-energy modes of the three-field quadratic tetrahedron at $\kappa/\mu = 10^4$, on the affine element. Script: `prototype/softmode.jl`.

Four consequences follow.

The constant-pressure element carries three modes that are exactly zero-energy on one element, not merely soft. Whether they persist after assembly is the question of §9.4.

A P_1 -discontinuous pressure removes them without an added energy term and leaves the six rigid-body modes. That is an element-level fact: §9.4 and §9.5 show that the same pair locks and is unstable once assembled.

Enriching the pressure further changes nothing: the rank saturates at four, because four is all the displacement space can produce. Enriching the pressure beyond P_1^{disc} adds unknowns and leaves the rank unchanged.

The count is a property of the affine element. On a curved element, with the midside nodes displaced so that $\det \mathbf{J}$ ranges over $[0.55, 1.19]$, the isoparametric displacement space is not polynomial in $\mathbf{x}$, only the dilatation of its four volumetric directions remains free of deviatoric strain, $\dim \ker \mathbf{K}^{\text{dev}} = 7$, and the three constant-pressure modes are soft rather than exactly zero: the softest nonzero mode of the constant-pressure element lies one to two orders of magnitude below the softest mode of the linear-pressure element on the same curved element. A four-point rule samples a curved element at four points and reports the affine counts there as well; the count uses a 27-point rule. The assembled behavior of §9.4 rests on soft rather than exact modes.

Remark 6. This identifies the mechanism. Inf-sup stability is a global property of a mesh sequence and is invisible on one element; the rank argument above says which modes exist, not whether the constant in the stability estimate persists under refinement; §9.5 measures it.

9.3 A necessary condition from degree-of-freedom counts

Whether the constant in the stability estimate persists under refinement is a property of a mesh sequence, and a necessary condition for it follows from counting.

Write n_u for the number of free displacement degrees of freedom and n_p for the pressure degrees of freedom. The Schur complement of the mixed system is $\mathbf{S} = \mathbf{G}\mathbf{K}^{-1}\mathbf{G}^\top$, an $n_p \times n_p$ matrix of rank at most n_u . If $n_p > n_u$ then $\mathbf{S}$ is singular by counting, the pressure space carries modes no displacement can resist, and no inf-sup constant exists — regardless of the values those constants would otherwise take.

Table 5 counts them on a Freudenthal-subdivided cube.

Pair	N	n_u	n_p	n_p/n_u
P_2/P_0	12	36,501	10,368	0.284
P_2/P_1^{disc}	12	36,501	41,472	1.136
$P_2 \oplus b_{\text{int}}/P_1^{\text{disc}}$	12	67,605	41,472	0.613
$P_2 \oplus b_{\text{int}} \oplus b_{\text{face}}/P_1^{\text{disc}}$	12	127,221	41,472	0.326

Table 5: Degree-of-freedom counts on the unit cube, displacement fixed on the boundary. Script: `prototype/inf-sup.jl`.

For P_2/P_1^{disc} the ratio exceeds one, and its limit is computable. The quadratic nodes of this mesh family fill the $(2N+1)^3$ grid, so with the whole boundary fixed $n_u = 3(2N-1)^3 \rightarrow 24N^3$, while $n_p = 4 \cdot 6N^3 = 24N^3$ regardless of boundary data. The ratio tends to one: the pair is critically determined in the limit, which excludes a uniform inf-sup constant.

The finite- N ratio depends on the boundary condition: with the whole boundary fixed, $n_u = 3(2N - 1)^3$ and the ratio exceeds one at every N ; with one face fixed, $n_u = 6N(2N + 1)^2$ and it is below one at every N . Both tend to one.

Proposition 1. *On a tetrahedral mesh family with at least six tetrahedra per vertex asymptotically — the Freudenthal family among them — and the boundary fixed, P_2/P_1^{disc} is not inf-sup stable, and refinement does not restore stability.*

The restriction to the mesh family is necessary. On a general tetrahedral mesh with n_V vertices, n_E edges and n_T tetrahedra, Euler’s relation for a simply connected three-dimensional mesh gives $n_E \approx n_V + n_T$, so $n_u \approx 3(2n_V + n_T)$, $n_p = 4n_T$, and $n_p/n_u \rightarrow 4/(3(2n_V/n_T + 1))$, which exceeds one only when $n_T/n_V > 6$. Below that the count refutes nothing, and on other mesh families the pair can be stable: Zhang [30] shows P_2/P_1^{disc} stable on Powell–Sabin tetrahedral meshes once their known spurious pressure modes are filtered. The general three-dimensional result is Boffi et al. [6, Remark 8.6.2]: with a P_k displacement, a discontinuous pressure is stable on general meshes only at degree $k - 3$, so for $k = 2$ nothing is admissible without enrichment.

Adding one interior bubble per element supplies $3n_{\text{el}}$ further displacement degrees of freedom and moves the ratio to 0.613; adding the four face bubbles as well — the enrichment §4.1 calls for — supplies $3n_f$ more, one vector unknown per interior face, and moves it to 0.326. Both satisfy $n_p < n_u$. The enrichment also fixes the number of global unknowns. The interior bubble condenses; the face bubbles are shared and do not. The global displacement count of the full enrichment is therefore $36,501 + 59,616 = 96,117$, which is 2.6 times the 36,501 of P_2/P_0 — also the nodal count of the composite tetrahedron on the same TETRA10 mesh — and 2.5 times the 38,698 global unknowns of Taylor–Hood, whose 2,197 vertex pressures cannot be eliminated. The 127,221 in Table 5 is the dimension of the space, not the number of global unknowns.

The count has precedents. n_p/n_u is the constraint count of Nagtegaal et al. [20], whose subject — locking of simplicial elements in the fully plastic range — is this note’s, and its reciprocal $\varrho = n_u/n_p$ is the constraint ratio of Hughes [16, §4.3.7], with $\varrho = 3$ the optimum in three dimensions. At $N = 12$: $\varrho = 3.52$ for P_2/P_0 , 0.88 for P_2/P_1^{disc} , 1.63 with the interior bubble, and 3.07 for the full enrichment. The stable pair has the optimal value, and the isochoric fraction $1 - n_p/n_u$ of §9.4 is $1 - 1/\varrho$, with $2/3$ the optimum.

Whether that limiting degeneracy is the decisive objection on the meshes tested is a separate question, which §9.4 answers.

9.4 Assembled behavior

The element count of §9.2 is a local statement. An element-local mode persists after assembly only if its trace is compatible with every neighbor across every shared face, so the global count can be smaller than $3n_{\text{el}}$, and it is what the argument against a constant pressure rests on.

Table 6 assembles $\mathbf{K}^{\text{dev}} + \kappa \mathbf{G}^T \mathbf{M}^{-1} \mathbf{G}$ over meshes under two boundary conditions: the whole boundary fixed, and the base $z = 0$ alone fixed. The second is closer to the conditions of a body confined on part of its boundary and free elsewhere, under which soft modes are reported.

The locking measure is again a rank. A displacement mode has volumetric energy exactly when it lies outside $\ker \mathbf{G}$, so as $\kappa \rightarrow \infty$ the reachable deformations collapse onto that kernel and $\dim \ker \mathbf{G} = n_u - \text{rank } \mathbf{G}$ is the discretely isochoric subspace. A large fraction means the element still deforms; a collapse toward zero is volumetric locking.

The fraction is compared with a bound. Since $\text{rank } \mathbf{G} \leq n_p$, it cannot fall below $1 - n_p/n_u$. A pair that attains that bound has $\text{rank } \mathbf{G} = n_p$: every pressure mode is resisted by a distinct displacement direction. A fraction above the bound means that $n_p - \text{rank } \mathbf{G}$ pressure modes are unresisted; a pair that cannot reach it, because $n_p > n_u$, is over-determined. The difference $n_p - \text{rank } \mathbf{G}$ counts the pressure modes no displacement resists, which with the whole boundary fixed should be exactly one, the hydrostatic mode; Table 7 carries it as a column.

boundary	Pair	N	n_u	n_p	$\text{rank } \mathbf{G}$	fraction	bound
whole	P_2/P_0	3	375	162	158	0.579	0.568
		4	1029	384	380	0.631	0.627
	P_2/P_1^{disc}	3	375	648	363	0.032	—
		4	1029	1536	957	0.070	—
	$\oplus b_{\text{int}}$	3	861	648	644	0.252	0.247
		4	2181	1536	1532	0.298	0.296
	$\oplus b_{\text{face}}$	3	1185	648	646	0.455	0.453
		4	3045	1536	1534	0.496	0.496
	$\oplus b_{\text{int}}, b_{\text{face}}$	3	1671	648	647	0.613	0.612
		4	4197	1536	1535	0.634	0.634
	$P_2 \oplus b_{\text{face}}/P_0$	3	1185	162	161	0.864	0.863
		4	3045	384	383	0.874	0.874
base	P_2/P_0	3	882	162	162	0.816	0.816
		4	1944	384	384	0.802	0.802
	P_2/P_1^{disc}	3	882	648	522	0.408	0.265
		4	1944	1536	1216	0.374	0.210
	$\oplus b_{\text{int}}$	3	1368	648	648	0.526	0.526
		4	3096	1536	1536	0.504	0.504
	$\oplus b_{\text{face}}$	3	1962	648	648	0.670	0.670
		4	4440	1536	1536	0.654	0.654
	$\oplus b_{\text{int}}, b_{\text{face}}$	3	2448	648	648	0.735	0.735
		4	5592	1536	1536	0.725	0.725
	$P_2 \oplus b_{\text{face}}/P_0$	3	1962	162	162	0.917	0.917
		4	4440	384	384	0.914	0.914

Table 6: Assembled count at $\kappa/\mu = 10^4$. The bound is $1 - n_p/n_u$ and is undefined where $n_p > n_u$. Assembled zero-energy modes were zero in every row. Script: `prototype/locking.jl`.

Three findings follow.

The constant-pressure soft modes do not persist after assembly. No spurious mode is present under either boundary condition at any mesh. Nor are they present as low-energy modes: the smallest eigenvalue relative to the deviatoric scale is 2.39×10^{-3} for P_0 against 2.74×10^{-3} for P_1^{disc} on the same mesh with the base alone fixed, and both are the long-wavelength modes of that block. The three modes per element of §9.2 exist on each element and are element-local: their traces do not match across faces.

What this test does not detect is the mode Foulk III et al. [14] report, which is remedied

there by a penalty and therefore has energy, small and not zero. Those are the displacement modes in $\ker \mathbf{G}_{P_0}$ whose divergence is nonzero pointwise, which carry shear energy alone in the element; whether they are spurious is the coercivity question of §6, which §9.7 measures, and not a count of zeros. The composite tetrahedron of Foulk III et al. [14] is built on a subdivided-linear displacement space rather than P_2 , which is a different kernel and is not modeled here; the soft mode of that element may be specific to it. Nothing above contradicts their observation; it does not reproduce it on a quadratic tetrahedron.

The unenriched pair locks. With the boundary fixed, P_2/P_1^{disc} retains 3–7% of its deformation modes as κ grows, against 58–63% for P_0 ; with the base alone fixed, 37–41% against 80%. Under either boundary condition it is the pair that over-constrains, and locking is the more serious failure because the element returns a stiff response and no spurious mode indicates the error.

The interior bubble raises the rank of $\mathbf{G}$ over a linear pressure. As Remark 7 states, the interior bubble adds nothing to the rank over P_0 . Against the linear pressure modes its columns are not zero: it raises $\text{rank } \mathbf{G}$ from 957 to 1532 at $N = 4$, and the interior-only pair has an isochoric fraction of 0.252–0.298. The full enrichment reaches 0.61–0.63, equal to P_0 ’s to two digits and without P_0 ’s element-level degeneracy; it attains the bound $1 - n_p/n_u$ to three digits and the interior-only pair to two, so every pressure mode is resisted by a distinct displacement direction; and for the full pair $n_p - \text{rank } \mathbf{G} = 1$ at every mesh, the hydrostatic mode alone, where P_2/P_0 and the interior-only pair carry four: the enrichment raises $\text{rank } \mathbf{G}$ to $n_p - 1$.

9.5 Stability, measured against a validated control

The count of §9.3 is a necessary condition; Table 7 shows that it is not sufficient.

Two properties of the computation allow meshes to $N = 8$. With $\mathbf{K} = \mathbf{P}^\top \mathbf{L} \mathbf{L}^\top \mathbf{P}$ from a sparse Cholesky, $q^\top \mathbf{G} \mathbf{K}^{-1} \mathbf{G}^\top q = \|\mathbf{L}^{-1}(\mathbf{G}^\top q)_{[p]}\|^2$, so β_h is the smallest nonzero singular value of $\mathbf{W} = \mathbf{L}^{-1} \mathbf{G}^\top \mathbf{M}^{-1/2}$ — never forming $\mathbf{G} \mathbf{K}^{-1} \mathbf{G}^\top$, and so never squaring the condition number of the quantity measured. And the null dimension is fixed independently as $n_p - \text{rank } \mathbf{G}$, so the cut is taken by index rather than by selecting a threshold. The gap is unambiguous: $O(10^{-15})$ below it against $O(10^{-2})$ above.

The sweep carries a positive control, the Taylor–Hood pair, whose stability in three dimensions is established [6]: a valid test must report it stable. A negative control alone cannot detect a test that reports every pair unstable; the negative control, P_1/P_0 , gives $\beta_h = 0.333, 0.229, 0.165, 0.129$ at $N = 2$ to 5 with 45, 138, 303, 558 unresisted pressure modes, the alternating-sign spurious pressure mode of that pair (script `prototype/infsup.jl`).

The control is reported stable: β_h rises and settles on 0.2216, flat to three digits over the last three refinements, which validates the reading of the remaining rows.

Every row of Table 7 is what Boffi et al. [6] predict. What the sweep establishes is that the scripts are sound and that the theory applies to this mesh family; what it adds to the theory is the constants. The rows are ordered as the construction of Boffi et al. [6, Example 8.7.2] adds degrees of freedom.

Pair	$h = 1/2$	$1/4$	$1/6$	$1/8$	null	local rate
P_2/P_1 continuous (control)	0.1734	0.2186	0.2214	0.2216	1	−0.003
P_2/P_0	0.1001	0.0755	0.0544	0.0418	4	0.97
$\oplus b_{\text{int}}$	0.1015	0.0756	0.0544	—	4	0.84
$\oplus b_{\text{face}}$	0.4671	0.4228	0.4060	0.3967	1	0.076
P_2/P_1^{disc}	0.1132	0.0424	0.0277	—	$n_p - n_u$	1.04
$\oplus b_{\text{int}}$	0.0782	0.0564	0.0406	—	4	0.83
$\oplus b_{\text{face}}$	0.0938	0.0701	0.0497	—	2	0.88
$\oplus b_{\text{int}}, b_{\text{face}}$	0.2875	0.2962	0.2968	—	1	−0.003

Table 7: β_h over a Freudenthal mesh sequence, whole boundary fixed, ordered as the construction of Boffi et al. [6, Example 8.7.2] runs. “null” is $n_p - \text{rank } \mathbf{G}$, the pressure modes no displacement resists, which must be one; “local rate” is $\ln(\beta_N/\beta_{N+1})/\ln((N+1)/N)$ at the last refinement, with β_N the constant at refinement N , constant for a genuine power law and shrinking toward zero for a sequence approaching a limit. Script: `prototype/beta.jl`.

P_2 does not control a constant pressure in three dimensions. P_2/P_0 decays at a local rate approaching one, although $n_p/n_u = 0.284$, and carries four pressure null modes at every mesh where one is allowed: the hydrostatic mode and three exact spurious modes that no quadratic displacement resists. That is the statement of Boffi et al. [6, §8.7] quoted in §4.1, and the three modes are the pressure modes that face degrees of freedom would resist.

The interior bubble cannot supply it. $P_2 \oplus b_{\text{int}}/P_0$ reproduces P_2/P_0 to three digits with the same four null modes, for the reason in Remark 7. Over P_1^{disc} the interior bubble changes β_h from 0.0424 to 0.0564 at $h = 1/4$ and from 0.0277 to 0.0406 at $h = 1/6$, and the rate falls from 1.04 to 0.83, but the sequence decays and four null modes remain. This pair has $n_p/n_u = 0.613$ (Table 5) and is unstable.

Face bubbles control the constant pressure. $P_2 \oplus b_{\text{face}}/P_0$ drops the null dimension to one at every mesh and its β_h approaches from above with a local rate that falls from 0.16 to 0.08 over the sequence, below the constant rate of a power law and above the rate −0.003 of the control. The seven refinements measured, of which Table 7 shows four, cannot separate a limit near 0.35 from a slow decay on their own, and the row is reported as consistent with stable; its stability rests on Boffi et al. [6, Example 8.7.2], whose proof constructs, from the face degrees of freedom, a displacement that matches the divergence of any given field in the mean on every face. Over P_1^{disc} the face bubbles alone leave one spurious linear pressure mode per mesh and β_h decays: controlling the linear modes is the function of the interior bubble.

Both together are the three-dimensional Crouzeix–Raviart pair, and it is stable. β_h rises monotonically to 0.2968, flat to three digits over the last three refinements, with a local rate of −0.003 — the same rate as the control’s, and at a higher constant than the control’s 0.2216. The pressure null space is one dimensional at every mesh. Each part of the enrichment controls one part of the pressure space, and neither controls the other’s.

Remark 7 (The distinct functions of the two enrichments). The interior bubble has zero pairing with a constant pressure: it vanishes on the whole element boundary, so $\int_{\Omega_e} \text{Div } \mathbf{b}_{\text{int}} = 0$ and its columns of $\mathbf{G}$ against a P_0 mode are identically zero. It pairs only with the linear pressure modes, which is why it changes P_2/P_1^{disc} and changes the P_2/P_0 constant by at most 0.0014 (the $P_2 \oplus b_{\text{int}}/P_0$ row of Table 7, with the same null dimension): upward, because enlarging U_h cannot lower β_h , and by at most 0.0014, because the bubble's columns of $\mathbf{G}$ against P_0 are zero. `checks.jl` reads those columns directly and confirms it. The face bubbles are what supply a displacement response to the pressure jump across a face, which is the mode a discontinuous pressure needs resisted.

P_2/P_0 and $P_2 \oplus b_{\text{int}}/P_1^{\text{disc}}$ have $n_p/n_u = 0.284$ and 0.613 and are unstable: a ratio above one excludes a pair, and a ratio below one does not admit it.

Remark 8 (Mesh dependence). These results are for the Freudenthal subdivision, a structured family, and structured meshes are where pair counterexamples characteristically arise: the alternating-sign spurious pressure mode of P_1/P_0 has its largest amplitude on them. A pair can degenerate on a structured mesh and remain well behaved on a general one. On this family, P_2/P_0 , P_2/P_1^{disc} and $P_2 \oplus b_{\text{int}}/P_1^{\text{disc}}$ decay; confirming that as a property of the pairs rather than of the mesh needs a sequence of unstructured meshes, which this study does not have. §9.6 gives one measurement in the opposite sense: on the non-affine images of this mesh the P_2/P_0 constant is smaller. The affirmative result needs no such confirmation, since Boffi et al. [6, Example 8.7.2] establish it on any shape-regular family; what the sweep adds is the constant.

9.6 Stability on a deformed configuration

Every measurement above is at $\mathbf{F} = \mathbf{I}$. At a finite deformation the pairing is $\int \varpi \text{div}_{\mathbf{x}} \boldsymbol{\psi} \, dV$ (§4.1), the divergence pairing on the deformed configuration with the pressure scaled by $1/J$, and whether β_h stays bounded below under a large non-affine deformation is what the regime of §1 requires. Auricchio et al. [4] show, on a two-dimensional finite-strain problem with a loading parameter, that mixed elements stable in linear elasticity (the MINI triangle, a mixed-enhanced quadrilateral and the quadrilateral Q_2/P_1^{disc}) fail to reproduce the range of the parameter over which the continuum problem is stable: the discrete problem becomes unstable where the continuum one is stable, or stays stable where it is not; and Auricchio et al. [5] conclude that for such pairs in finite elasticity the causes are identified but a general remedy is still missing. This section measures the constant of the pair on deformed configurations for that reason. The operators (49) are assembled on the isoparametric image of the quadratic mesh under a prescribed map, with the derivatives, the divergence and the volume taken on the image and the pressure basis on the reference cell as the element uses it, and with the boundary conditions of the reference cube; the scaling by $1/J$ changes β_h by at most the ratio of extreme Jacobians and is not applied. Five maps of the unit cube, about its center: a twist about the vertical axis by 90° and by 180° over the height, isochoric and non-affine, with shear strains near 2 at the corners of the latter; a parabolic shear $x_1 = X_1 + \frac{1}{2}X_3^2$, isochoric; a radial inflation with J from 1 to 2.3; and an affine compression to $J = \frac{1}{2}$ as a control, since an affine map changes β_h only through the element aspect ratio. Table 8 gives the result at $N = 5$ with the local rate over the last refinement. Script: `prototype/deformed.jl`.

Three findings follow.

Table 8: β_h at $N = 5$ on deformed configurations, whole boundary fixed, with the local rate over $N = 4 \rightarrow 5$ in parentheses for the Crouzeix–Raviart pair; the null dimension is one on every row for the three left-hand pairs. The P_2/P_0 column gives its null dimension and β_h : its three exact spurious modes exist only on the affine mesh.

map	J	P_2/P_1 cont.	$P_2 \oplus b_{\text{face}}/P_0$	$P_2 \oplus b_{\text{int}} \oplus b_{\text{face}}/P_1^{\text{disc}}$	P_2/P_0 (null, β_h)
reference	1	0.221	0.413	0.297 (−0.003)	4, 0.063
twist 90°	0.99–1.01	0.189	0.385	0.204 (0.18)	1, 0.027
twist 180°	0.98–1.05	0.137	0.345	0.116 (0.42)	1, 0.051
parabolic shear	1	0.176	0.336	0.183 (0.05)	1, 0.005
inflation	1.0–2.3	0.192	0.368	0.276 (0.01)	1, 0.004
compression	0.5	0.192	0.333	0.213 (0.00)	4, 0.038

β_h is not uniform in the deformation. It decreases with the local shear of the map for every pair, and most for the Crouzeix–Raviart pair: from 0.297 to 0.204 under the 90° twist, 0.116 under the 180° twist and 0.183 under the parabolic shear, against 0.221 $\rightarrow$ 0.189, 0.137, 0.176 for Taylor–Hood and 0.413 $\rightarrow$ 0.385, 0.345, 0.336 for the face-bubbled constant-pressure pair. The inflation, which changes J by a factor of 2.3 but shears little, leaves the constant at 0.276. The affine compression gives 0.213, flat to four digits from $N = 3$: the aspect-ratio effect alone. The dependence is on the shear of the deformation and not on its volume change.

At a fixed deformation the sequence in h is consistent with a positive limit. The local rate of the Crouzeix–Raviart pair falls with N on every non-affine map: 0.28, 0.22, 0.18 over $N = 2$ to 5 for the 90° twist, 0.67, 0.53, 0.42 for the 180° twist, and 0.05 for the parabolic shear (through $N = 5$). A power-law decay would hold its rate. The 180° row has not flattened by $N = 5$, and its shear strain of about 2 at the corners is beyond what a plastic zone of the problems listed in §12 reaches; the finding is that the constant degrades with the strain, not that it vanishes. At $N = 6$ the pair gives 0.1975, 0.1088 and 0.182 on the three non-affine maps, local rates of 0.17, 0.35 and 0.04, each below the rate before it.

The exact spurious modes of P_2/P_0 belong to the affine mesh. Under every non-affine map the pressure null dimension of P_2/P_0 falls from four to one, the three modes are resisted, and β_h of P_2/P_0 is smaller on every non-affine image, 0.004 to 0.05, than on the reference mesh, 0.063.

9.7 Coercivity under a plastic tangent

The claim of §6 is measured as the constant it names. With $\mathbf{Z}$ a basis of $\ker \mathbf{G}_h$,

$$\lambda_{\min} = \min_q \frac{q^\top \mathbf{Z}^\top \mathbf{K}^{\text{ep}} \mathbf{Z} q}{q^\top \mathbf{Z}^\top \mathbf{K}^{H^1} \mathbf{Z} q}, \quad \mathbf{K}_{ij}^{\text{ep}} = \frac{1}{2\mu} \int \boldsymbol{\epsilon}_i : \mathbb{C}^{\text{ep}} : \boldsymbol{\epsilon}_j \, dV, \quad (53)$$

is the smallest generalized eigenvalue of the deviatoric form under the tangent (47) against the H^1 seminorm, restricted to the discretely isochoric subspace. It is dimensionless, it is taken on the exact kernel rather than through a penalty, and it is the constant of the second hypothesis; (48) says what to expect of it, at least $(1 - \eta)/2$ for every pair when $\eta < 1$ and a decay to

Table 9: Lowest mode of the deviatoric form on $\ker \mathbf{G}_h$ under the tangent (47) at $\eta = 1$, Freudenthal cube with the whole boundary fixed: the eigenvalue λ_1 of (53) and the dilatation r_v of its mode, (54), for a shear and a uniaxial flow direction. The elastic lowest mode has $\lambda_1 = 0.500$ and $r_v = 0$ on every row, as the identity behind (48) requires. At $\eta = 0.9$ every λ_1 exceeds the lower bound 0.05 of (48) and the dilatations agree with those below to two digits. Script: `prototype/plastic.jl`.

Pair	N	$\mathbf{n} = \text{sym}(\mathbf{e}_1 \otimes \mathbf{e}_2)/\sqrt{2}$		$\mathbf{n} = \text{dev}(\mathbf{e}_3 \otimes \mathbf{e}_3)/\ \cdot\ $	
		λ_1	r_v	λ_1	r_v
P_2/P_0	3	0.109	0.065	0.117	0.208
	4	0.067	0.034	0.106	0.217
	5	0.047	0.029	0.101	0.220
$P_2 \oplus b_{\text{face}}/P_0$	3	0.044	0.044	0.082	0.274
	4	0.026	0.028	0.074	0.268
	5	0.017	0.019	0.071	0.261
$P_2 \oplus b_{\text{int}} \oplus b_{\text{face}}/P_1^{\text{disc}}$	3	0.074	0.036	0.147	0.018
	4	0.049	0.019	0.139	0.0074
	5	0.032	0.0066	0.134	0.0042
P_2/P_1 continuous (control)	3	0.056	0.051	0.053	0.861
	4	0.030	0.029	0.026	0.929
	5	0.019	0.020	0.016	0.956

zero for every pair when $\eta = 1$. The quantity that carries the claim is the dilatation of the eigenmodes,

$$r_v(\mathbf{v}) := \frac{\int (\text{Div } \mathbf{v})^2 dV}{\int |\text{Grad } \mathbf{v}|^2 dV}, \quad (54)$$

zero for an isochoric field. The flow direction is uniform, $\mathbf{n} = \text{sym}(\mathbf{e}_1 \otimes \mathbf{e}_2)/\sqrt{2}$ or $\mathbf{n} = \text{dev}(\mathbf{e}_3 \otimes \mathbf{e}_3)/\|\cdot\|$, which is the most unfavorable choice: the continuum operator then admits shear bands of any width, a decaying eigenvalue is expected of every pair, and only the dilatation of the mode says whether the band is physical. The tangent is taken at $\eta = 1$ and at $\eta = 0.9$, the boundary is fixed everywhere, and each pair is measured elastically ($\eta = 0$) on the same kernel. Script: `prototype/plastic.jl`.

The scripts are checked by an identity. By the identity behind (48), the elastic form on any kernel satisfies $\lambda = \frac{1}{2} + \frac{1}{6}r_v$, so its lowest mode has $r_v = 0$ and $\lambda_1 = 0.5$ exactly. Every row of the sweep reproduces this; at $N = 2$ for P_2/P_0 , where the kernel has no isochoric member, $\lambda_1 = 0.5024$ with $r_v = 0.0146$, as the identity requires. The identity is also the check on the energy metric of the scripts.

The constant does not separate pairs, as (48) predicts. Under the shear direction λ_1 decays for every pair, the control included, at rates between $h^{1.5}$ and h^2 : the discrete spaces resolve narrower shear bands, which the continuum operator admits with zero energy. Under the uniaxial direction no shear band has $\epsilon \propto \mathbf{n}$, and λ_1 converges for the three element-local pairs, to 0.10 for P_2/P_0 , 0.071 for $P_2 \oplus b_{\text{face}}/P_0$ and 0.13 for the Crouzeix–Raviart pair, while

Table 10: Spurious modes among the twenty lowest of the plastic form on $\ker \mathbf{G}_h$ ($r_v > 0.1$) in a confined plastic zone with a varying flow direction, at $N = 3, 4, 5$.

field	boundary	P_2/P_0	$P_2 \oplus b_{\text{face}}/P_0$	$P_2 \oplus b_{\text{int}} \oplus b_{\text{face}}/P_1^{\text{disc}}$	P_2/P_1 cont.
indentation	whole fixed	16, 15, 7	18, 17, 5	4, 0, 0	20, 18, 17
indentation	base fixed	11, 1, 0	13, 5, 0	0, 0, 0	14, 7, 5
twist	whole fixed	14, 18, 15	18, 13, 8	1, 4, 0	20, 20, 20
twist	base fixed	5, 2, 0	13, 7, 4	0, 0, 0	12, 13, 12

it decays for Taylor–Hood. At $\eta = 0.9$ every λ_1 lies above the lower bound 0.05 of (48), the smallest being the control’s 0.067.

The dilatation of the soft modes separates them, in the direction claimed. Under the uniaxial direction the lowest modes of both constant-pressure pairs carry $r_v = 0.20$ – 0.27 at every refinement, with no decay. These are the modes of §6: the element assigns them no energy for a dilatation the continuum resists with κ , and they are the softest modes the element has. The lowest modes of the Crouzeix–Raviart pair carry 0.040, 0.018, 0.0074 and 0.0042 over $N = 2$ to 5, decaying as h^2 , the rate at which Boffi et al. [6, Prop. 5.1.3] makes $\ker \mathbf{G}_h$ approximate the isochoric space; its soft modes are physical. Under the shear direction the lowest modes of every pair are shear bands whose dilatation decays, the Crouzeix–Raviart pair’s three to four times smaller than the others’ at $N = 5$. Counted over the twenty lowest modes, with a mode called spurious when $r_v > 0.1$ (the stable pair’s soft modes lie below 0.02 from $N = 4$ on, the spurious ones at 0.2–0.4): under the uniaxial direction at $N = 5$ all twenty of each constant-pressure pair’s and of Taylor–Hood’s are spurious, and none of the Crouzeix–Raviart pair’s at any N ; under the shear direction the twenty lowest of every element-local pair are bands at $N = 5$. The spurious modes of a constant pressure are therefore present in both directions but are the lowest modes only where the flow admits no band; the next paragraph therefore counts them among the twenty lowest in a zone with a varying flow direction.

A confined zone with a varying flow direction. The flow direction and the plastic zone are taken from a prescribed displacement field: $\mathbf{n} = \text{dev } \boldsymbol{\epsilon}(\mathbf{u}) / \|\text{dev } \boldsymbol{\epsilon}(\mathbf{u})\|$ at each quadrature point and $\eta = 1$ where $|\text{dev } \boldsymbol{\epsilon}(\mathbf{u})|$ exceeds 0.35 of its peak, $\eta = 0$ elsewhere. Two fields: an indentation of the top face, $u_3 = -\exp(-d^2/w^2) X_3$ with d the distance from the vertical axis and $w = \frac{1}{4}$, whose zone is 51% of the volume with elastic surroundings; and the displacement of the 90° twist of §9.6, whose zone is 97%. Both are run with the whole boundary fixed and with the base alone fixed. Table 10 counts the spurious modes among the twenty lowest. Script: `prototype/plastic_zone.jl`.

With the whole boundary fixed the constant-pressure pairs retain between five and fifteen spurious modes among their twenty softest at $N = 5$, and under the twist the softest mode of P_2/P_0 is spurious at every N up to 4 ($r_v = 0.22$ – 0.31) and its third at $N = 5$ (0.20). The Crouzeix–Raviart pair has none from $N = 4$ on under the indentation and none at $N = 5$ under the twist, and none at any N with the base alone fixed. Taylor–Hood keeps seventeen to twenty, and under the twist its softest mode’s dilatation grows with refinement, 0.29, 0.37, 0.66, 0.74. With the base alone fixed the softest modes of every pair are long-wavelength motions of

the block and the counts fall for the constant-pressure pairs too, to zero at $N = 5$ for P_2/P_0 and four for the face-bubbled pair, so the spurious modes are then above the twentieth; they are not absent, as the whole-boundary rows show. Whether a nonlinear solve excites them is the open measurement of §12.

Taylor–Hood is the worst pair on this measure. Its pressure is continuous, so its kernel is divergence-free only against continuous linear test functions, and under the uniaxial direction its lowest modes are dilatation almost entirely: $r_v = 0.86, 0.93, 0.96$ at $N = 3, 4, 5$, with λ_1 decaying to 0.016. Those modes have no counterpart in the continuum, which resists them with κ . Taylor–Hood is therefore an alternative for elasticity only (§13).

The composite tetrahedron on the same measure. The composite tetrahedron of Albany-LCM, the element §10 compares against, is measured with the same scripts (`prototype/composite_tet.jl`) as Albany builds it: ten nodes, a displacement that is linear on each of twelve subtetrahedra, the gradient of that displacement projected onto the linear functions of the element for the constitutive update, and a five-point rule of degree three. Two volumetric variants exist in Albany: J replaced by its element mean (the option `Weighted Volume Average J`), whose constraint as $\kappa \rightarrow \infty$ is a constant pressure on the projected dilatation, and J at each quadrature point, whose constraint at the five points is a linear discontinuous pressure. With the averaged J , under the uniaxial flow direction, the three lowest modes carry $r_v = 0.22$ to 0.25 at $N = 3, 4, 5$ with no decay, all twenty lowest modes are spurious from $N = 4$ on, and in the confined zones of Table 10 with the whole boundary fixed the count at $N = 5$ is 17 under the indentation and 19 under the twist, against 7 and 15 for P_2/P_0 . With J at each point the discretely isochoric subspace has dimension 0, 0, 2 and 16 at $N = 2$ to 5, against 72 for P_2/P_1^{disc} at $N = 4$: the element locks, and too few modes exist to count. The average removes the locking and admits the spurious dilatational modes of a constant pressure, in larger numbers than P_2/P_0 . The measurement is linearized, like every other in this section, and models no stabilization of the element’s own zero-energy modes, because the Albany runs of §10 use none.

9.8 Quadrature rule and basis

Two implementation choices set the number of operations per element and its explicit stable step, and neither is decided by the stability results: the quadrature rule, and the basis that carries the enriched space. Both are measured on one element. Scripts: `prototype/quadrature.jl` and `prototype/basis.jl`.

Quadrature. The scripts integrate the enriched element to degree 6 so that the linear measurements are exact. For the element the criterion is weaker: a rule is admissible when the condensed stiffness $\mathbf{K}_e = \mathbf{K}^{\text{dev}} + \kappa \mathbf{G}^T \mathbf{M}^{-1} \mathbf{G}$ has exactly six zero-energy modes, and its effect is the softening it introduces, measured as the extreme generalized eigenvalues of $\mathbf{K}_e(\text{rule})$ against $\mathbf{K}_e(\text{exact})$ on the complement of the rigid modes, which are 1 for an exact rule. Table 11 gives both for the rules considered, with a degree-11 rule of 216 points as the reference. The two rules supplied by `ReferenceFiniteElements` (RFE), the Julia package that provides the reference elements and their quadrature, four and five points, leave 21 zero modes and are inadmissible. The eight-point conical rule (a tensor-product Gauss rule mapped

Table 11: Quadrature rules on the enriched element: zero-energy modes of the condensed stiffness (six is correct) and of the deviatoric stiffness (ten: six rigid modes and the four volumetric directions of §9.2), and the extreme softening factors against a 216-point rule, on the reference tetrahedron and on a Kuhn tetrahedron of the Freudenthal mesh. $\kappa/\mu = 10^4$. Script: `prototype/quadrature.jl`.

rule	points	reference			Kuhn		
		zero($\mathbf{K}_e$)	min	max	zero($\mathbf{K}_e$)	min	max
RFE degree 2	4	21	0	1.30	21	0	1.30
RFE degree 3	5	21	< 0	1.13	21	< 0	1.12
conical degree 3	8	6	0.10	1.14	6	0.03	1.21
Keast degree 5	14	6	1.00	1.35	6	1.00	1.35
conical degree 5	27	6	0.91	1.00	6	0.88	1.00
conical degree 7	64	6	1.00	1.00	6	1.00	1.00

onto the tetrahedron) is rank sufficient and softens the element to 0.03–0.10 of its stiffness in the worst direction, which is also inadmissible. Two degree-5 rules are admissible: the 14-point symmetric rule of Keast [17], which integrates every monomial through degree 5 to 4×10^{-16} and over-stiffens the interior bubble’s own block by 35% since its gradient-squared is of degree 6, and the 27-point conical rule, which under-stiffens the same block by 9–12%. Assembled, both give the Crouzeix–Raviart pair’s constant within 1% (0.293 and 0.298 against 0.297 at $N = 5$), a one-dimensional pressure null space, no spurious mode among the twenty lowest under the uniaxial plastic tangent, and the same dilatation of the lowest mode (0.0073, 0.0074 and 0.0074). The 14-point rule satisfies the criteria — six zero-energy modes, positive weights, 4.6 times fewer points than the 64-point degree-7 rule — and is used, with a stiffer interior bubble.

Basis. The space $P_2 \oplus b_{\text{int}} \oplus b_{\text{face}}$ can be carried by the hierarchical basis of the scripts, ten Lagrange functions and five bubbles, or by the nodal basis with unit value at one of fifteen nodes, the **TETRA15** arrangement: vertices, edge midpoints, face centroids, centroid. The space is the same, so every stiffness and stability result is the same; what differs is conditioning, lumping, and the coefficients of a rigid translation. Table 12 gives the measurements at $\kappa/\mu = 2$. The nodal basis conditions the consistent mass eight times better (94 against 778) and the stiffness better (559 against 797). Under HRZ lumping the hierarchical basis loses 51% of the momentum of a rigid translation unless the normalization runs over the Lagrange functions alone (§7); the nodal basis conserves it without qualification, and its lumped masses are positive on every element measured. The explicit stable step from the element bound, as the Courant number $c_p \Delta t/h$, is 0.476 for TET4, 0.197 for TET10 under HRZ, 0.145 for the enriched element in the nodal basis under HRZ, and 0.114 in the hierarchical basis under Lagrange-only HRZ, with the same ordering on the distorted and Kuhn elements. The nodal basis is used: every unknown is attached to a node, so no unknown of a new kind is introduced; its mass lumps correctly by the standard construction; and it takes a 20–27% larger stable step.

Table 12: Basis choice on the reference tetrahedron, $\kappa/\mu = 2$, $\mu = \rho = 1$: condition numbers of the condensed stiffness on the complement of the rigid modes and of the consistent mass, the smallest lumped mass over ρV , the momentum of a unit translation under the lumped mass over ρV , and the Courant number $c_p \Delta t/h$ from the element bound. Script: `prototype/basis.jl`.

basis, lumping	cond $\mathbf{K}_e$	cond $\mathbf{M}$	min $m/\rho V$	momentum	Courant
hierarchical, consistent	797	778	—	—	0.092
hierarchical, HRZ over all 15	—	—	0.014	0.486	0.097
hierarchical, HRZ over Lagrange 10	—	—	0.028	1.000	0.114
nodal, consistent	559	94	—	—	0.092
nodal, HRZ	—	—	0.021	1.000	0.145
nodal, row sum	—	—	0.020	1.000	—
TET10, HRZ	—	—	0.028	1.000	0.197
TET10, row sum	—	—	−0.050	1.000	—
TET4, row sum	—	—	0.250	1.000	0.476

10 Nonlinear results: Cook’s membrane

Every measurement of §9 is on the linearized operator. This section measures the element in nonlinear solves, against the composite tetrahedron of Albany-LCM on the same meshes, on Cook’s membrane: a tapered panel clamped on one edge and sheared on the opposite one [10], in its nearly incompressible form [27] and in its elastoplastic form, both in three dimensions.

Problem. The membrane is the trapezoid with corners $(0, 0)$, $(48, 44)$, $(48, 60)$ and $(0, 44)$ in the x - y plane, of thickness 10 in z , with both z faces free. The face $x = 0$ is clamped in all three components and the face $x = 48$ carries a uniform shear traction q in y , applied in ten equal load steps. The material is the Simo–Hughes J_2 model (Remark 5) in both codes. Two cases: the elastic one, with Young’s modulus $E = 240.565$, Poisson’s ratio $\nu = 0.4999$ ($\kappa/\mu = 5000$), a yield stress that is never reached and $q = 6.25$, the plane-strain load $F = 100$ per unit thickness divided by the height 16 of the loaded edge; and the plastic one, with the material of Simo and Armero [27], $E = 206.9$, $\nu = 0.29$, yield stress $\sigma_y = 0.45$, linear isotropic hardening modulus $K = 0.12924$, and $q = 0.14$. Their load $q = 0.1125$ leaves the three-dimensional membrane with free faces almost elastic (tip displacement 0.21 with no difference between elements); the quadratic tetrahedron without projection collapses between $q = 0.16$ and $q = 0.20$, and Albany’s Newton iteration on the composite tetrahedron stops converging at $q = 0.158$ on the middle mesh and $q = 0.137$ on the finest, so the load is 0.14. The measured quantity is the mean of the displacement u_y over the nodes of the loaded face at full load.

Meshes and elements. Three TETRA10 meshes with element size $h = 8, 4$ and 2 (327, 1860 and 14 475 elements; 647, 3208 and 22 059 nodes) are generated by Cubit and used by both codes; the TETRA15 meshes are their conversions (§4.3; 1721, 9127 and 66 813 nodes). Six elements are compared. In Carina: TET10 with the volumetric response evaluated at each quadrature point (pointwise); TET15 pointwise; TET15-P1, the element of this note, with the linear projection; and TET15-P0, the same displacement element with the constant projection. In Albany-LCM: TET10 pointwise, which serves as the cross-code control; and the composite

Table 13: Cook’s membrane, mean u_y of the loaded face at full load. Rows: element size h with the element count. Elastic case: $E = 240.565$, $\nu = 0.4999$, $q = 6.25$. Plastic case: $E = 206.9$, $\nu = 0.29$, $\sigma_y = 0.45$, $K = 0.12924$, $q = 0.14$. Driver: `benchmark/tet15-p1/cook/run.jl`; records in `results.tsv`.

h (elements)	TET10 pointwise	TET15 pointwise	TET15-P1	TET15-P0	TET10 Albany	composite tet. (Albany)
elastic						
8 (327)	8.1960	8.2235	8.4122	8.5720	8.1960	8.4596
4 (1860)	8.3241	8.3490	8.4551	8.5184	8.3241	8.4643
2 (14475)	8.4203	8.4300	8.4716	8.4996	8.4203	8.4772
plastic						
8 (327)	0.2689	0.2700	0.2706	0.2752	0.2689	0.2716
4 (1860)	0.2710	0.2717	0.2721	0.2744	0.2710	0.2725
2 (14475)	0.2720	0.2723	0.2724	0.2734	0.2720	0.2728

tetrahedron with the volume-averaged J and pressure. The displacement unknowns are 1941, 9624 and 66 177 on the TETRA10 meshes and 5163, 27 381 and 200 439 on the TETRA15 meshes. Both codes converge each load step to a residual norm of 10^{-6} absolute or 10^{-8} relative with a direct linear solver; both halve a load step whose Newton iteration fails and grow it back after successful steps.

Cross-code control. The pointwise TET10 gives the same tip displacement in Carina and in Albany to six digits on every mesh and in both cases (8.19605 elastic at $h = 8$; 8.4202946 and 8.4202943 at $h = 2$; 0.27204 plastic at $h = 2$), although the two codes take the volumetric stress as $\kappa(J - 1)$ and $\kappa(J - 1/J)/2$: at these volumetric strains the two laws agree to the digits shown. The comparison between the other elements rests on this check.

Tip displacement. Table 13 gives the results. In the elastic case the pointwise elements converge from below, 8.196, 8.324, 8.420 for TET10 and 8.224, 8.349, 8.430 for TET15, and TET15-P0 from above, 8.572, 8.518, 8.500. TET15-P1, 8.412, 8.455, 8.472, and the composite tetrahedron, 8.460, 8.464, 8.477, are within 0.6% of each other on every mesh and change least with refinement; the limit lies between 8.477 and 8.500. On the coarsest mesh the pointwise TET10 is 3% below the limit and TET15-P1 0.8% below it. In the plastic case every element converges to about 0.273, with the spread between elements shrinking from 2.3% at $h = 8$ to 0.5% at $h = 2$, and the projected and composite elements closest to the limit on every mesh; at $\nu = 0.29$ and this load the plastic zone does not lock the quadratic elements. On the finest mesh the elastic runs took 1099 s for TET10 and 3214 to 3639 s for the TET15 variants in Carina on 16 threads of one host, and 26 496 and 28 867 s in Albany with a serial direct solver; at $\nu = 0.4999$ the residual grows with κ times the square of the volumetric strain of a Newton step, so every element needs load steps of a fraction of a percent of the load in both codes.

Pressure fields. The pressure $p = -\text{tr } \sigma/3$, with σ the Cauchy stress, is read at the quadrature points of every element at full load and drawn on the face $z = 10$ of the deformed membrane (Figures 1–3). Inside each element the field drawn is the polynomial fitted to the




Figure 1: Elastic case, $h = 2$: pressure $p = -\text{tr } \sigma / 3$ on the face $z = 10$ of the deformed membrane at full load, for the six elements. One color scale for all panels, symmetric about zero and bounded by the 99th percentile of $|p|$ over TET15-P1, TET15-P0 and the composite tetrahedron; values beyond the scale saturate. Inside each element the pressure drawn is the polynomial fitted to the quadrature-point values (text).




Figure 2: Plastic case, $h = 2$: pressure on the face $z = 10$ at full load, as in Figure 1.




Figure 3: Pressure on the face $z = 10$ under mesh refinement, $h = 8, 4, 2$ from left to right in each panel, with the color scale of the case. Top: TET15-P1, elastic (left) and plastic (right). Bottom left: TET10 pointwise, elastic. Bottom right: the composite tetrahedron, plastic.

quadrature-point values in the least-squares sense with the weights of the rule: quadratic for the fourteen-point rule of TET15, which reproduces the linear and the constant projection to rounding error and the pointwise TET15 pressure, of higher degree, to 0.11 of its largest value; linear for the four-point rule of TET10, which interpolates; and the single value Albany writes per composite tetrahedron. In the elastic case the pointwise TET10 and TET15 give a pressure that changes sign inside elements over the whole membrane, on every mesh, with a range on the face of -226 to 211 and -512 to 1010 at $h = 2$. TET15-P1 is smooth except in the layer of elements along the clamped edge, with a range of -43 to 61 ; TET15-P0 and the composite tetrahedron are smooth and piecewise constant, with ranges of -7.5 to 22 and -9.4 to 35 . The clamped edge is a singular line of the continuum solution: where a clamped face meets a free face the stress is unbounded, and it grows faster as $\nu \rightarrow 1/2$. The jumps of TET15-P1 between neighboring elements confirm that the discontinuities are the element's error at that line and nothing else: at every h the six largest jumps on the face lie within one element layer of the clamped edge, the largest $|p|$ of TET15-P1 on the face grows from 22 to 35 to 61 over the three meshes as an unbounded field requires, and away from that layer and from the loaded corner the median jump falls from 3.9% to 0.35% of the color limit from $h = 8$ to $h = 2$, faster than h , whereas for the two piecewise constant elements it is 12 to 15% of the color limit at $h = 2$, in proportion to h . In the plastic case no element oscillates, the six fields agree away from the upper corner of the clamped edge, and the largest pressure on the face at $h = 2$ is 0.715 for TET10, 1.15 for TET15 pointwise, 0.519 for TET15-P1, 0.292 for TET15-P0 and 0.407 for the composite tetrahedron. Scripts: `benchmark/tet15-p1/cook/extract_pressure.jl` and `render_pressure.py`.

What the benchmark establishes. The element of this note does not lock at $\nu = 0.4999$ on any of the three meshes, agrees in the tip displacement with the composite tetrahedron within 0.6% and in the limit of refinement, and gives a pressure that is smooth wherever the continuum pressure is, on a mesh where the quadratic tetrahedron without projection gives a pressure that oscillates in sign everywhere. Its unknowns are 3.0 times those of the composite tetrahedron on the same mesh, the interior bubble not being condensed in the implementation, and the run times above are those of one implementation with a direct solver; a comparison at equal accuracy is not made. The equivalent plastic strain is not measured here: at $q = 0.14$ the plastic zone of the membrane does not lock any of the quadratic elements, so the membrane does not separate them on that quantity.

11 Explicit dynamics: the Taylor bar impact

This section measures the element in explicit dynamics with plastic strains above 3 , against the composite tetrahedron of Sierra/SM, the code of Foulk III et al. [14], on the same meshes. The quantity measured is the final radius of the impact face, for which a converged reference value exists; the equivalent plastic strain is not measured.

Problem. A copper bar of length 32.4 mm and radius 3.2 mm, with its axis along z , strikes a rigid frictionless wall at $z = 0$ with velocity 227 m/s [14, 24, 29]. The material has density $\rho = 8930$ kg/m³, Young's modulus $E = 117$ GPa, Poisson's ratio $\nu = 0.35$ ($\kappa/\mu = 3.0$; the plastic flow is isochoric, so the response in the plastic zone is nearly incompressible), yield

stress 400 MPa and a linear isotropic hardening modulus of 100 MPa; the Simo–Hughes J_2 model (Remark 5) in Carina and the finite-deformation J_2 model of the Sierra material library, with the same constants. The wall is the condition $u_z = 0$ on the impact face, which is in contact at $t = 0$ and remains in contact while the bar moves toward the wall. The equations are integrated by central differences with the row-sum lumped mass of §7 to $t = 80\mu\text{s}$, without artificial bulk viscosity. The measured quantities are the radius of the impact face, the largest distance of its nodes from the axis, and the length of the bar, the largest z of its free end, at every microsecond. Foulk III et al. [14] give 7.22 mm for the final radius of the eight-node hexahedron with constant pressure (Q1/P0) on their finest mesh, $h = 0.047$ mm, which is the reference here; Simo [24] gives 7.12 mm and Zienkiewicz et al. [31] 7.11 mm on coarser meshes.

Meshes and elements. Four levels, $h = 1.5, 0.75, 0.38$ and 0.19 mm, of the full bar are meshed by Cubit with four-node tetrahedra, the Cubit size being adjusted until the element count is within 3% of that of Foulk III et al. [14, Figure 14]: 3522, 25 029, 190 727 and 1 521 212 elements, with 5601, 36 666, 267 150 and 2 083 320 nodes as TETRA10 meshes and 16 587, 113 331, 845 573 and 6 672 877 as TETRA15 meshes. Each mesh is smoothed by the energetic mesh smoothing of Norma, which minimizes a strain energy of the mesh over the positions of its nodes with every boundary node held on its surface; this raises the smallest shape quality $q = 12(3V)^{2/3}/\sum_e l_e^2$, with V the volume and l_e the six edge lengths, from 0.58–0.65 to 0.69–0.71 at every level. The edge midpoints are added with those on the lateral surface projected onto the cylinder (TETRA10, the mesh of the composite tetrahedron), and the face and interior nodes are added by the conversion of §4.3 (TETRA15). Three elements are compared: TET15-P1 and TET15-P0 in Carina, and the composite tetrahedron in Sierra/SM with its stabilization parameter $\alpha = 0.1$ and the stabilization exponent 0 of Foulk III et al. [14]. The displacement unknowns of TET15 are 49 761, 339 993 and 2 536 719 on the first three levels; those of the composite tetrahedron are 109 998, 801 450 and 6 249 960 on the second to fourth.

Stable time step. The step of central differences is stable when $\Delta t \leq 2/\sqrt{\lambda_{\max}}$, with $\lambda_{\max}$ the largest eigenvalue of $M_L^{-1}K$, M_L the lumped mass and K the tangent stiffness. The runs of Table 14 set the step to a Courant number of 0.25 times an element estimate, the smallest distance between two nodes of an element divided by the dilatational wave speed c_p , recomputed every ten steps on the current configuration. That estimate is not a bound: $\lambda_{\max}$ computed by power iteration on the TET15-P1 run at $h = 0.75$ mm gives a ratio of the critical step to the element estimate of 1.14 at $1\mu\text{s}$, 1.02 at $10\mu\text{s}$, 0.69 at $20\mu\text{s}$, 0.42 at $28\mu\text{s}$, 0.31 at $40\mu\text{s}$ and 0.29 at 60 and $80\mu\text{s}$. At a Courant number of 0.5 the run stopped with a non-finite residual at $28\mu\text{s}$; at 0.25 it reached $80\mu\text{s}$ with a step 0.82 times the critical one at $40\mu\text{s}$; TET15-P0 at $h = 0.38$ mm stopped at $78.7\mu\text{s}$, and its radius at $78\mu\text{s}$ is reported, the radius being constant within 0.06% after $50\mu\text{s}$ for every element (Figure 4). Carina also offers the step $0.8 \cdot 2/\sqrt{\lambda_{\max}}$, with $\lambda_{\max}$ computed every 200 steps and the element estimate scaled by the ratio of the two in between; on the same run it takes 43 863 steps against about 54 000 at the Courant number 0.25, and the final radius differs by 0.0007 mm.

Accuracy. Table 14 and Figure 4. The composite tetrahedron converges from above: its differences between successive levels are 0.0367 and 0.0127 mm, a ratio of 2.89, which corresponds to an order of 1.53 in h ; the extrapolation of its three levels by that ratio gives 7.2259 mm. TET15-P1 is at 7.2225 and 7.2255 mm at $h = 0.75$ and 0.38 , within 0.05% of that value, and

Table 14: Taylor bar, final radius of the impact face and final length of the bar in mm at $80\mu\text{s}$, per element and mesh level h in mm with its element count. Reference: 7.22 for the converged Q1/P0 hexahedron [14]. TET15-P0 at $h = 0.38$ stopped at $78.7\mu\text{s}$; its values at $78\mu\text{s}$. The composite tetrahedron at $h = 1.5$ ended at $41.6\mu\text{s}$. Driver: `benchmark/tet15-p1/taylor/run.jl`; histories in `benchmark/tet15-p1/taylor/data`.

h	elements	TET15-P1		TET15-P0		composite tet.	
		radius	length	radius	length	radius	length
1.5	3522	7.2399	21.4250	7.4415	21.5507	—	—
0.75	25 029	7.2225	21.4206	7.4747	21.5329	7.2820	21.4760
0.38	190 727	7.2255	21.4202	7.3785	21.4763	7.2453	21.4395
0.19	1 521 212	—	—	—	—	7.2326	21.4304




Figure 4: Taylor bar. Left: final radius of the impact face against the number of elements; the open square is TET15-P0 at $78\mu\text{s}$. Right: radius of the impact face against time for TET15-P1 and TET15-P0 at $h = 0.38$ mm and the composite tetrahedron at $h = 0.38$ and 0.19 mm. Dashed gray: the converged Q1/P0 hexahedron of Foulk III et al. [14]. Script: `benchmark/tet15-p1/taylor/plot.py`.

Table 15: Taylor bar at $h = 0.75$ mm (25 029 elements), both codes on the same two-socket AMD EPYC 9634, each run alone. One-core figures from runs to $5\mu\text{s}$. Core- μs : wall time times the number of cores. Memory of Sierra/SM: the largest of its processes.

	TET15-P1 (Carina)	composite tet. (Sierra/SM)
per element and step, 1 core (μs)	6.27	4.47
per element and step, 16 cores (core- μs)	15.5	5.09
steps to $80\mu\text{s}$	52 100	44 260
time stepping to $80\mu\text{s}$ on 16 cores (s)	1265	352
largest resident memory (GB)	1.63	0.27
final radius (mm)	7.2225	7.2820

at 7.2399 on the coarsest mesh, where the composite tetrahedron did not reach $80\mu\text{s}$. The composite tetrahedron is 0.78%, 0.27% and 0.09% from the extrapolated value at $h = 0.75$, 0.38 and 0.19. Both elements approach the reference within 0.1% in the limit. TET15-P0 is 3.4% above the extrapolated value at $h = 0.75$ and 2.1% at $h = 0.38$; it differs from TET15-P1 only in the constant pressure.

Operations. The two codes were run on the same machine, a two-socket AMD EPYC 9634 with 168 cores, one run at a time with no other job, at $h = 0.75$ mm (Table 15). Per element and time step on one core, TET15-P1 takes 1.40 times the time of the composite tetrahedron; per element degree of freedom, 45 against 30, it takes 0.93 times as long. On 16 cores the ratio per element and step is 3.0, because Carina’s time per step decreases by a factor of 6.5 from 1 to 16 threads, a parallel efficiency of 0.40, against 0.88 for Sierra/SM on 16 processes; the difference is in the implementation of the time step in Carina, not in the element. The time step of TET15-P1 at $t = 0$ is 3.86×10^{-9} s under the element estimate at the Courant number 0.25, and its critical step is 1.77×10^{-8} s; that of the composite tetrahedron is 2.44×10^{-8} s.

Comparison at equal accuracy. TET15-P1 at $h = 0.75$ mm is closer to the extrapolated radius (0.05%) than the composite tetrahedron at $h = 0.19$ mm (0.09%). On the machine above, each run alone, the first takes 1303 s on 16 cores, 2.1×10^4 core-seconds, and the second 10 983 s on 128 processes, 1.41×10^6 core-seconds: a factor of 67 in core-seconds and of 8.4 in wall time, on 25 029 elements against 1 521 212. At equal mesh the composite tetrahedron takes 3.6 times less time to $80\mu\text{s}$ on 16 cores, with an error 16 times larger.

The same run on GPUs. The explicit integration of TET15-P1 runs on a GPU through the same kernels, with the stable step recomputed on the device. At $h = 0.75$ mm with the global stable step, the radius and the length of every output frame of the run are the same numbers to every written digit on the CPU, on an AMD RX 7600 and on NVIDIA L4, V100 and A100 cards (the last three on one mesh, the first two on another smoothing of the same Cubit mesh, which changes the radius by 2.8×10^{-7} relative), and so are the 221 stable-step estimates. Table 16 gives the time per step from the kernel timings and the wall time of the run. Three defects were found and corrected in making these runs: the stable-step recomputation read the block sizes and the material properties through accessors defined for host arrays only, so it had never run on a device; the temporaries of the explicit step (about 0.65 MB per step at 340 000 unknowns) are released only when the garbage collector runs, which host pressure did

Table 16: TET15-P1, split form, explicit Taylor bar: time per step of a run from the kernel timings (median of 20, deformed state with plastic strains to 1.0; includes the stable-step estimates at their intervals), in microseconds per element-step, at $h = 0.75$ and 0.38 mm, the ratio of the general form’s time to the split form’s at $h = 0.75$ mm, and the wall time of the run to $80\mu\text{s}$ at $h = 0.75$ mm. Carina 9e4bfaf; one process alone on each machine. Scripts: `benchmark/tet15-p1/taylor/kernel_timing.jl`, `timing_table.py`; records in `data/`.

device	$h = 0.75$	$h = 0.38$	general/split	run (s)
AMD Ryzen 9 9900X, 12 threads	0.662	—	1.61	839
2×AMD EPYC 9634, 32 threads	0.700	0.772	1.36 [†]	1283 [†]
2×AMD EPYC 9634, 64 threads	0.836	0.652	—	—
AMD RX 7600	0.496	—	1.44	513
NVIDIA L4	0.365	0.322	1.60	406
NVIDIA V100	0.325	0.250	2.05	404
NVIDIA A100	0.169	0.119	1.84	227

[†]at 16 threads.

not trigger, so the run exhausted the 8 GB of the RX 7600 after about 5 000 steps, and the collector is now run every 100 steps on a device; and the J_2 tangent of the material library built its pull-back in mutable buffers that CUDA places on the device heap, which is never freed, so the matrix-free solves of the implicit tests failed on the L4 after some thousands of tangent calls (the buffers are now immutable, and the tangent is 5% faster on every device). The eigenvalue estimate of the stable step takes 25 to 80 explicit steps’ time and runs every 200 steps, so it is 13% of the stepping time on a CPU with 16 threads and 25 to 40% on the GPUs; the ratio it corrects varies by 1% per 200 steps late in the run, so the interval can be raised on a GPU.

Pressure field. Figure 5 shows the pressure $p = -\frac{1}{3} \text{tr } \sigma$, positive in compression, on the impact face at $t = 12\mu\text{s}$, the time of Figure 12 of Foulk III et al. [14], for both elements on the same meshes. The composite tetrahedron has one value per element, the element mean of the stress that Sierra/SM writes; for TET15-P1 the pressure in each element is the quadratic fitted to its fourteen quadrature-point values, as in §10. The largest pressure on the impact end is 0.33, 0.48, 1.37 and 1.59 GPa for the composite tetrahedron at $h = 1.5, 0.75, 0.38$ and 0.19 mm, and 1.67, 1.58, 1.56 and 1.54 GPa for TET15-P1, between 3% below and 5.2% above the composite tetrahedron’s value on its finest mesh at every level; on that mesh the two differ by 3%. On the two coarser meshes the pressure of the composite tetrahedron varies from element to element without a maximum on the axis; TET15-P1 gives the maximum on the axis on every mesh. The radius of the impact face at $12\mu\text{s}$ is 5.4944, 5.4957, 5.4750 and 5.4661 mm for the composite tetrahedron and 5.4517, 5.4507, 5.4533 and 5.4570 mm for TET15-P1. Foulk III et al. [14] plot the L^2 projection of the composite tetrahedron’s pressure onto a continuous field, and the range of their color scale, -1.2 to 0.32 GPa, indicates the mean stress $\frac{1}{3} \text{tr } \sigma = -p$. Scripts: `benchmark/tet15-p1/taylor/pressure.jl`, `render_pressure.py` and `compose_pressure.py`.

What the benchmark establishes. In explicit dynamics with plastic strains above 3, TET15-P1 gives the final radius of the Taylor bar within 0.05% of the converged value on a mesh of 25 029 elements, where the composite tetrahedron needs 1.5 million elements to reach




Figure 5: Taylor bar, pressure $p = -\frac{1}{3} \text{tr } \boldsymbol{\sigma}$ on the impact face at $t = 12\mu\text{s}$, seen along the axis, on the same meshes for both elements, with the element count of each level and the node count of each mesh (**TETRA10** for the composite tetrahedron, **TETRA15** for TET15-P1); one color scale, values beyond it saturate. Top: composite tetrahedron (Sierra/SM, $\alpha = 0.1$, exponent 0), one value per element. Bottom: TET15-P1, the quadratic fitted in each element.

0.09%, and the constant-pressure variant on the same displacement space is 2.1 to 3.4% from it. On 3522 elements its largest pressure on the impact face at $12\mu\text{s}$ is 5.2% above that of the composite tetrahedron on 1.5 million, and on 1.5 million elements the two differ by 3%. Its time per element and step is 1.40 times that of the composite tetrahedron on one core. The element estimate of the stable step used by Carina is not a bound on distorted elements; the eigenvalue estimate replaces it. The equivalent plastic strain was not written in these runs and is not compared.

12 Summary of the evidence and open measurements

The formulation rests on six claims. For each, the section that establishes it and the number that carries it are given; the measurements that remain are listed at the end.

The pair is inf-sup stable in three dimensions. §9.5: the Crouzeix–Raviart pair has $\beta_h = 0.2968$ at a local rate of -0.003 , above the Taylor–Hood control’s 0.2216, with a one-dimensional pressure null space at every mesh; the constant-pressure pair with face bubbles is stable with a constant decreasing toward a limit near 0.35, and the five other element-local pairs decay. The face bubbles are the necessary part of the enrichment; the interior bubble alone is not sufficient (Remark 7). Uniformity in κ/μ follows from stability of the incompressible limit [6, §5.5.2, §8.12.1, Remark 8.12.1]. On deformed configurations (§9.6) the constant decreases with the shear of the deformation, to 0.204 at a 90° twist and 0.116 at 180° , and by 7% under a volume change of 2.3, and at each fixed deformation the sequence in h is consistent with a positive limit; the constant is not uniform in the deformation.

No spurious soft modes. In elasticity the second hypothesis of Brezzi’s theorem holds for every conforming pair (§6). The three exact zero-energy modes per element of a constant pressure (§9.2) do not persist after assembly (§9.4). Under a plastic tangent (§9.7) the coercivity constant is bounded below by (48) for every pair, and the dilatation of the soft modes separates them: 0.20–0.27 without decay for the constant-pressure pairs, the composite tetrahedron with its averaged J among them at 0.22–0.25, h^2 decay for the Crouzeix–Raviart pair, and in a confined zone with a varying flow direction no spurious mode among the twenty softest of the Crouzeix–Raviart pair at $N = 5$ against five to nineteen for the constant-pressure elements.

No volumetric locking. §9.4: the discretely isochoric fraction $\dim \ker \mathbf{G}/n_u$ is 0.61–0.63 for the Crouzeix–Raviart pair, at the bound $1 - n_p/n_u$, against 0.03–0.07 for P_2/P_1^{disc} , which locks. §10: on Cook’s membrane at $\nu = 0.4999$ the element agrees with the composite tetrahedron in the tip displacement within 0.6% on meshes of 327 to 14 475 elements, and its pressure is smooth away from the singular clamped edge, where the quadratic tetrahedron without projection oscillates in sign inside every element.

Accuracy in explicit dynamics. §11: on the Taylor bar the final radius is within 0.05% of the converged value on 25 029 elements, against 0.09% for the composite tetrahedron on 1 521 212, at $1/67$ of its core-seconds on the same machine; the constant-pressure variant is 2.1 to 3.4% from it.

Element-local elimination. §4.2: the pressure is discontinuous, so its elimination is element by element and the tangent has the sparsity of the displacement mesh; the face bubbles add global displacement unknowns, 2.6 times those of P_2/P_0 at $N = 12$ (§9.3).

Quadrature and basis. §9.8: the 14-point degree-5 rule keeps six zero-energy modes and the assembled constant within 1%; the nodal basis has a positive row-sum lumped mass that conserves the momentum of a rigid translation, and a Courant number of 0.145 under HRZ lumping against 0.476 for TET4.

Open measurements. Three remain. An unstructured mesh sequence for the negative stability results, which are established here on the Freudenthal family only (Remark 8). The smoothness of the equivalent plastic strain in problems with localized plastic flow, the benchmark set on which the composite tetrahedron performs poorly [14]: punch indentation, large-rotation twist of a bar, and a weak shock in a material of high κ/μ , measured with the implemented Simo–Hughes instance (Remark 5) against the composite tetrahedron on the same meshes; Cook’s membrane at the load of §10 does not separate the elements on that quantity, and the Taylor bar runs of §11 did not write it. The comparison at equal accuracy against the composite tetrahedron has been made on one problem, the Taylor bar (§11); a problem that the quadratic element does not resolve on a coarser mesh would give a different result.

13 Limitations

This section states what the formulation requires, what it does not establish, and what its parts are.

The number of unknowns. The face bubbles of the Crouzeix–Raviart pair are shared between neighbors and do not condense: at $N = 12$ the enriched space carries 96,117 global displacement degrees of freedom against 36,501 for P_2/P_0 and for the composite tetrahedron on the same mesh, a factor of 2.6, and 2.5 times Taylor–Hood’s 38,698 (§9.3); the implementation does not condense the interior bubble, and on the Cook meshes of §10 the factor is 3.0. On the Taylor bar (§11) the accuracy on a coarser mesh more than compensates for that factor: the final radius within 0.05% on 25 029 elements, against 1.5 million elements for the composite tetrahedron at 0.09%. On other problems it is not established; for elasticity that is the comparison with Taylor–Hood, whose pressure is global and whose displacement space is 2.5 times smaller, and for plasticity Taylor–Hood’s softest modes are dilatation almost entirely (§9.7), so the comparison there is with the composite tetrahedron.

Operations per element. With the 14-point rule of §9.8 the enriched element evaluates the constitutive law at 14 points on 45 displacement degrees of freedom, against 4 or 5 points on 30 for TET10 and one on 12 for TET4. On the Taylor bar the time per element and explicit step on one core is 1.40 times that of the composite tetrahedron in Sierra/SM (§11).

The explicit stable step. Row-sum lumping in the nodal basis gives positive masses and conserves momentum, and the stable step is 3.3 times smaller than TET4’s on the same element, of which a factor of 2.4 is the quadratic element and 1.4 the enrichment (§9.8). On the

Taylor bar at equal accuracy the smaller element count outweighs the smaller step (§11). On a deformed mesh an element estimate from node distances does not bound the critical step; a step of 0.8 times $2/\sqrt{\lambda_{\max}}$, with $\lambda_{\max}$ the largest eigenvalue of $M_L^{-1}K$, stays below it, at a cost of 13% of the stepping time on a CPU and 25 to 40% on the GPUs with the estimate every 200 steps (Table 16); a longer interval on a GPU is not yet tested. If weak shocks are run with the implicit integrator the factor does not apply.

The material. The element places no requirement on the material: the split form of §4.4 calls it through the split of §3.5, which of the seven material models of Table 1 five have, and the general form of §3.6 calls any material at the modified deformation gradient (25), at 1.36 to 1.60 times the split form’s time per residual. What the general form requires is a positive volumetric tangent of the material over the states reached, and what it changes for a material whose isochoric response depends on J is that this response is evaluated at the projected volume. The tangent of the general form is assembled from fifteen forward-mode passes per element; the closed form (29) is not yet coded, and the time of the tangent and of the preconditioner kernels is not measured. For the split form, the action of the tangent on a vector takes 2.4 times the residual’s time on the RX 7600 (20.8 against 8.6 ms on the Taylor mesh at $h = 0.75$ mm).

The parts. Mixed volumetric treatment [26], logarithmic-strain plasticity [18] and displacement–pressure pairs satisfying (37) uniformly [6] are established, and the pair itself has been used in large-deformation hyperelasticity [8]. The contribution is their combination for finite-deformation plasticity: an element-local pair satisfying (37) uniformly, a projected volumetric strain, and a deviatoric law evaluated pointwise, which removes the stabilization that the composite-tetrahedron and variational-multiscale methods carry; together with the measurements that support it, of the dilatation of the soft modes under a plastic tangent (§9.7), of the inf-sup constant on deformed configurations (§9.6), and of the nonlinear response (§10, §11). The evidence is that of §9, §10 and §11; the smoothness of the equivalent plastic strain is the open measurement of §12. Hybrid high-order methods, with unknowns on element faces and element interiors coupled only through the faces, are locking-free and eliminate their interior unknowns element by element, and exist for small-deformation plasticity [1] and for finite plasticity in logarithmic strain [2]; they are the alternative, at the price of face unknowns that are not the nodal displacements of a standard element.

References

- [1] M. Abbas, A. Ern, and N. Pignet. A hybrid high-order method for incremental associative plasticity with small deformations. *Computer Methods in Applied Mechanics and Engineering*, 346:891–912, 2019.
- [2] Mickaël Abbas, Alexandre Ern, and Nicolas Pignet. A hybrid high-order method for finite elastoplastic deformations within a logarithmic strain framework. *International Journal for Numerical Methods in Engineering*, 120(3):303–327, 2019. doi: 10.1002/nme.6137.
- [3] A. H. Al-Mohy, N. J. Higham, and S. D. Relton. Computing the Fréchet derivative of the matrix logarithm and estimating the condition number. *SIAM Journal on Scientific Computing*, 35:C394–C410, 2013.

- [4] Ferdinando Auricchio, Lourenco Beirão da Veiga, Carlo Lovadina, and Alessandro Reali. A stability study of some mixed finite elements for large deformation elasticity problems. *Computer Methods in Applied Mechanics and Engineering*, 194(9–11):1075–1092, 2005. doi: 10.1016/j.cma.2004.06.014.
- [5] Ferdinando Auricchio, Lourenco Beirão da Veiga, Carlo Lovadina, Alessandro Reali, Robert L. Taylor, and Peter Wriggers. Approximation of incompressible large deformation elastic problems: some unresolved issues. *Computational Mechanics*, 52(5):1153–1167, 2013. doi: 10.1007/s00466-013-0869-0.
- [6] D. Boffi, F. Brezzi, and M. Fortin. *Mixed Finite Element Methods and Applications*, volume 44 of *Springer Series in Computational Mathematics*. Springer, 2013.
- [7] Franco Brezzi. On the existence, uniqueness and approximation of saddle-point problems arising from Lagrangian multipliers. *Revue Française d’Automatique, Informatique et Recherche Opérationnelle. Analyse Numérique*, 8(R2):129–151, 1974.
- [8] Éric Chamberland, André Fortin, and Michel Fortin. Comparison of the performance of some finite element discretizations for large deformation elasticity problems. *Computers & Structures*, 88(11–12):664–673, 2010. doi: 10.1016/j.compstruc.2010.02.007.
- [9] Dominique Chapelle and Klaus-Jürgen Bathe. The inf-sup test. *Computers & Structures*, 47(4–5):537–545, 1993.
- [10] Robert D. Cook. Improved two-dimensional finite element. *Journal of the Structural Division, ASCE*, 100(ST9):1851–1863, 1974.
- [11] M. Crouzeix and P.-A. Raviart. Conforming and nonconforming finite element methods for solving the stationary Stokes equations I. *Revue Française d’Automatique, Informatique et Recherche Opérationnelle*, 7:33–75, 1973.
- [12] Eduardo A. de Souza Neto, Djordje Perić, Martin Dutko, and D. Roger J. Owen. Design of simple low order finite elements for large strain analysis of nearly incompressible solids. *International Journal of Solids and Structures*, 33(20–22):3277–3296, 1996.
- [13] Michel Fortin. An analysis of the convergence of mixed finite element methods. *RAIRO. Analyse Numérique*, 11(4):341–354, 1977.
- [14] J. W. Foulk III, J. T. Ostien, B. Talamini, M. R. Tupek, N. K. Crane, A. Mota, and M. G. Veilleux. Extending a 10-node composite tetrahedral finite element for solid mechanics. *International Journal for Numerical Methods in Engineering*, 122:6939–6967, 2021.
- [15] E. Hinton, T. Rock, and O. C. Zienkiewicz. A note on mass lumping and related processes in the finite element method. *Earthquake Engineering and Structural Dynamics*, 4:245–249, 1976.
- [16] Thomas J. R. Hughes. *The Finite Element Method: Linear Static and Dynamic Finite Element Analysis*. Prentice-Hall, Englewood Cliffs, NJ, 1987.
- [17] Patrick Keast. Moderate-degree tetrahedral quadrature formulas. *Computer Methods in Applied Mechanics and Engineering*, 55(3):339–348, 1986.

- [18] C. Miehe, N. Apel, and M. Lambrecht. Anisotropic additive plasticity in the logarithmic strain space: modular kinematic formulation and implementation based on incremental minimization principles for standard materials. *Computer Methods in Applied Mechanics and Engineering*, 191:5383–5425, 2002.
- [19] A. Mota, W. Sun, J. T. Ostien, J. W. Foulk III, and K. N. Long. Lie-group interpolation and variational recovery for internal variables. *Computational Mechanics*, 52:1281–1299, 2013.
- [20] J. C. Nagtegaal, D. M. Parks, and J. R. Rice. On numerically accurate finite element solutions in the fully plastic range. *Computer Methods in Applied Mechanics and Engineering*, 4:153–177, 1974.
- [21] J. T. Ostien, J. W. Foulk III, A. Mota, and M. G. Veilleux. A 10-node composite tetrahedral finite element for solid mechanics. *International Journal for Numerical Methods in Engineering*, 107:1145–1170, 2016.
- [22] Stefan Sauter and Céline Torres. On the inf-sup stability of Crouzeix–Raviart Stokes elements in 3D. arXiv:2205.00062 [math.NA], 2022.
- [23] G. Scovazzi, B. Carnes, X. Zeng, and S. Rossi. A simple, stable, and accurate linear tetrahedral finite element for transient, nearly, and fully incompressible solid dynamics: a dynamic variational multiscale approach. *International Journal for Numerical Methods in Engineering*, 106:799–839, 2015.
- [24] J. C. Simo. Algorithms for static and dynamic multiplicative plasticity that preserve the classical return mapping schemes of the infinitesimal theory. *Computer Methods in Applied Mechanics and Engineering*, 99:61–112, 1992.
- [25] J. C. Simo and T. J. R. Hughes. *Computational Inelasticity*, volume 7 of *Interdisciplinary Applied Mathematics*. Springer, New York, 1998.
- [26] J. C. Simo, R. L. Taylor, and K. S. Pister. Variational and projection methods for the volume constraint in finite deformation elasto-plasticity. *Computer Methods in Applied Mechanics and Engineering*, 51:177–208, 1985.
- [27] Juan C. Simo and Francisco Armero. Geometrically non-linear enhanced strain mixed methods and the method of incompatible modes. *International Journal for Numerical Methods in Engineering*, 33(7):1413–1449, 1992.
- [28] P. Thoutireddy, J. F. Molinari, E. A. Repetto, and M. Ortiz. Tetrahedral composite finite elements. *International Journal for Numerical Methods in Engineering*, 53:1337–1351, 2002.
- [29] M. L. Wilkins and M. W. Guinan. Impact of cylinders on a rigid boundary. *Journal of Applied Physics*, 44(3):1200–1206, 1973. doi: 10.1063/1.1662328.
- [30] Shangyou Zhang. Quadratic divergence-free finite elements on Powell–Sabin tetrahedral grids. *Calcolo*, 48:211–244, 2011.

- [31] O. C. Zienkiewicz, J. Rojek, R. L. Taylor, and M. Pastor. Triangles and tetrahedra in explicit dynamic codes for solids. *International Journal for Numerical Methods in Engineering*, 43(3):565–583, 1998. doi: 10.1002/(SICI)1097-0207(19981015)43:3<565::AID-NME454>3.0.CO;2-9.